

Non-Intrusive load monitoring

RUNE A. HEICK 11061

Aarhus University, Department of Engineering

April 1, 2016

Abstract

Write an abstract Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.

Abbreviations

ECO Electricity Consumption and Occupancy. 29–31, 35, 42, 55

EMD empirical mode decomposition filling. 15

FHMM factorial hidden Markov models. 26, 32, 37, 41–43, 45, 46, 49–51, 54, 58

HMM hidden Markov models. 24–27, 32, 50, 55

IMF intrinsic mode functions. 15

NILL non-intrusive load leveling. 63

NILM non-intrusive load monitoring. 3, 6, 10, 11, 21, 22, 24–26, 31, 41, 43, 50, 53, 56, 57, 59–63

NILMTK NILM-Toolkit. 6, 10, 25, 27

NIST National Institute of Standards and Technology. 1

P-G Papoulis-Gerchberg. 13, 14, 17, 19

SSA Spatio-Temporal filling. 14

UDP user datagram protocol. 5

List of Figures

1.1	Conceptual model of smart-grid architecture. Source [?]	2
2.1	12 hour overview of house 10	7
2.2	Illustration of availability analysis	8
2.3	Quality of houses in SmartHG project	9
2.4	Activity of houses in SmartHG project	10
2.5	Usage of SmartHG houses	11
3.1	Gap quantity	15
3.2	Error recovery capability	15
3.3	Available samples for recovery	16
3.4	Sample comparison of the reconstruction methods	17
3.5	Stochastic effect on prediction.	18
3.6	Frequency comparison of the reconstruction methods	19
3.7	Jitter Power Illustration	19
3.8	Jitter Power comparison of the reconstruction methods	20
4.1	Appliance types. Source [1] make your own figure	22
4.2	Feature types. Source [1]	23
4.3	Factorial hidden Markov model illustration. Source [2]	27
4.4	Parson hidden Markov model structure. Source [3]	28
4.5	House energy distribution. Source [4]	30
4.6	Algorithm validation at different sample graduality.	32
4.7	Appliances accuracy score	33
4.8	Appliances F1 score	34
4.9	Algorithm validation at different error rate	35
4.10	Average recognition after recovery of the 3 methods	36
4.11	Recovery in FHMM algorithm	37
4.12	Recovery in Parson algorithm	38
4.13	Recovery in Weiss algorithm	39
5.1	Disaggregation phases	41
5.2	Frequency comparison of the reconstruction methods	42
5.3	FHMM disaggregation snippet	43
5.4	Artificially constructed main meter	44
5.5	Illustration of dataset creation by combining real and constructed data	46
5.6	Illustration of dataset creation by combining constructed and real data	47
5.7	Disaggregation of TV 1 from house 10	49
5.8	Purged and merged signal from TV 1	50
5.9	\mathbf{H}_c sets from the TV house dataset	52
5.10	Average score at different complexity	53

5.11	Completeness test of disaggregation of appliances in the TV house dataset	54
6.1	The city setup	56
6.2	The disaggregation process	57
6.3	Viewership in week one	58
7.1	TV event detection	62
7.2	NILL illustration. Inspired by [?]	63

List of Tables

5.1	Appliance disaggregation results of house 3,10 and 18 for the smartHG dataset . . .	43
5.2	Disaggregation of appliances on artificially constructed main meters	45
5.3	Disaggregation in real and constructed main meters combined	47
5.4	Disaggregation in constructed and real main meters combined	48
5.5	Trained in real vs. constructed data	48
5.6	TV House dataset complexity	52
6.1	Average viewership in a 6 week period	59

Contents

Abbreviations	ii
List of Figures	iii
List of Tables	iv
Contents	v
1 Introduction	1
1.1 Smart Grid	1
1.2 Non-Intrusive Load Monitoring	3
1.3 The Problem Statement	3
1.3.1 The Data Approach	4
2 Data Quality	5
2.1 Quality Criteria	6
2.1.1 Related Work	6
2.2 Quality In SmartHG Citizen Data	7
2.2.1 Completeness And Activity	7
2.2.2 Main Meter In Relation To Sub-meter consumption	10
2.3 Chapter Discussion	11
3 Gap Reconstruction	13
3.1 Gap Filling Methods	13
3.1.1 Papoulis-Gerchberg Algorithm	13
3.1.2 Wiener Filling Algorithm	14
3.1.3 Spatio-Temporal Filling Algorithm	14
3.1.4 Envelope Filling Algorithm	14
3.1.5 Empirical Mode Decomposition Filling Algorithm	15
3.2 Gaps In SmartHG Dataset	15
3.2.1 Gap Size	15
3.2.2 Post And Prior Knowledge	16
3.3 SmartHG Dataset Reconstruction	16
3.3.1 Sample Comparison	17
3.3.2 Frequency Comparison	19
3.3.3 Jitter Comparison	19
3.4 Chapter Discussion	20
4 Appliance Recognition	21
4.1 NILM Concepts And Challenges	21
4.1.1 NILM Features	22
4.1.2 Learning Strategy	23

4.1.3	NILM Challenges	24
4.2	Related Work	25
4.3	Recognition Methods	26
4.3.1	Factorial Hidden Markov Models	26
4.3.2	Parson	27
4.3.3	Weiss	29
4.4	The ECO Dataset	29
4.5	Validation of Methods	30
4.5.1	Sample rate experiment	31
4.5.2	Error Tolerance Experiment	35
4.5.3	Gap Filling Experiment	35
4.6	Chapter Discussion	39
5	Environment Influence	41
5.1	Challenges In The SmartHG Dataset	42
5.2	Background Consumption Influence On Disaggregation	42
5.2.1	Detection In An Environment With High Background Consumption	43
5.2.2	Detection In An Environment With No Background Consumption	44
5.2.3	Background Consumption Effect On Training And Validation	46
5.2.4	Improving Performance Using Norm Knowledge	49
5.3	Model Size And Completeness Influence	50
5.3.1	Test Set Creation	51
5.3.2	Model Complexity Test	53
5.3.3	Model Completeness Influence	54
5.4	Chapter Discussion	55
6	NILM As An Application	56
6.1	The TV Viewing Habits Case	56
6.2	Methodology	57
6.3	Results	58
6.4	Chapter Discussion	59
7	Discussion	60
7.1	Creating High Quality Measurements	60
7.2	Capability Of NILM	61
7.3	NILM Deployment	61
7.4	NILM Privacy Concerns	62
8	Conclusion	64
	Bibliography	65

In the past electricity production and consumption was localised to small communities. The communities often had one electricity utility company supplying the required energy. As the energy demand increased beyond single production facilities, the solution was to add more production facilities that could support in peak hours. This approach is costly, since such support facilities must be kept in a standby state, ready to deliver additional energy at a moments notice. As the electric grid expanded it began to interconnect the communities. The distribution of energy in the grid also became a challenging task. Energy must be delivered at the different consumers, at correct quantities, without overloading the grid. Furthermore must the impact errors and damages have on the grid be minimised.

1.1 Smart Grid

As sensor technology became cheaper and more reliable the electricity utility companies started to integrate sensors in the grid. These sensors delivered information about the performance of key sections of the grid. This information was used to improve the performance of the grid. The grid grew to be more global and less local. More types of energy producing devices, such as solar power and wind turbines, got connected to it. This created a need for more control of the grid and the energy production.

The National Institute of Standards and Technology (NIST) is tasked with the job of creating and maintaining the standards used for the American smart-grid. The concept behind the smart-grid is to create a electrical grid that enables two way communication. Besides electricity must information about the usage also flow between the consumers and the utility companies. Usage patterns must flow to the engineers tasked with controlling the grid, and control information must flow back to the consumers or control points in the grid. This can ensure the delivery of electricity more efficiently, reliably, and securely [?].

There are many different kind of stakeholders in today's smart-grid. In order to better understand the needs and relationships between the stakeholders have NIST split them in to seven domains. This helps to identify the different user interest in the grid, and their needs.




Figure 1.1: Conceptual model of smart-grid architecture. Source [?]]

Figure 1.1 illustrates the seven domains. The seven domains each have a set of functions that is characteristic for the domain [?]].

Generation: Generators of electricity. This is typically coal, oil or nuclear power plants or large-scale hydro generators. This can also include energy storage facilities that stores energy for later distribution.

Transmission: Transmission and transformers that hold the purpose of transporting the electricity over large distances.

Distribution: The units responsible of electricity distribution to and from the customers.

Customer: The consumers of electricity. They may also generate electricity and sell it to the grid.

Operator: The managers of the movement and production of electricity. They are the group that ensures that the energy is efficiently delivered where it is needed.

Market: The sale and purchase of electricity between the different instances of the grid. One utility company could purchase power from another to keep up with demand at peak hour.

Service Provider: The organisations providing services to the utility companies or the consumer. This could be home automation or information services.

1.2 Non-Intrusive Load Monitoring

The operators have the responsibility of controlling the flow of energy on the smart grid. This requires the knowledge of the current consumption and a prediction about the future power requirement. This prediction is generally based on statistical data collected in the grid. If the prediction of the future power requirements can be improved, it is possible to create a better and more efficient power distribution. To accomplish this is a method known as non-intrusive load monitoring (NILM) developed.

The concept of NILM is to use the consumption information collected at household level, and use machine learning techniques to make a qualified guess on what appliances in the household is responsible for the current consumption. By using the information from the active appliances more precise estimates of future consumption can be made.

In today's smart grid is all customers equipped with a smart-meter. The smart-meter measures the consumption of the household, and reports it back to the smart grid. This information can be used by the markets to automatically create billing information, or by the operator to regulate the power distribution. Since this information is used for distribution regulation it is sent in real-time [?]. It is this information NILM applications hope to use in order to create better predictions.

NILM can also create many opportunities for the service provider. Using NILM it is potentially possible to gesture about the usage of each appliance in a household. This can be used to better inform the resident about power saving opportunities. Studies have shown that detailed information about energy usage, makes the resident use less energy [5]. Other business opportunities, like selling the statistical information about the usage of specific appliances, also exists.

1.3 The Problem Statement

This report seeks to investigate some of the common approaches used for NILM today, in order to determine the capabilities of such approaches. It is assumed that such approaches are used in a modern smart grid setting. In order to answer this broad question, the following questions must be investigated.

- What quality of data can be expected from the smart-meters?
- How do errors and poor quality affect NILM applications?
- What can we detect using NILM and where are the limits?
- Can anything be done to improve the NILM technology?

If these questions are answered it creates a strong foundation for gesturing about the capability of NILM. The main motivation for this study is not to improve grid load predictions, but to get a better basis for validating different NILM related business opportunities for service providers.

1.3.1 The Data Approach

To investigate the questions in the problem statement is the SmarHG dataset used. The SmarHG dataset is a collection of smart-meter data collected from 25 residential households in Denmark. The data is not filtered or cleaned in any way prior usage. This means that the data closely resembles what would be collected in a smart grid. Furthermore does the dataset contain information about the usage pattern of selected devices. This information can be used to compare inferred usage patterns to the actual.

Data Quality 2

Various projects today is focused on gathering data and analysing it. The gathered data is used for obtaining behaviours, habits and properties of the observed objects. This is done by using powerful statistical leaning algorithms, that are able to deduce these properties from the data. This approach is called data driven development, since the success is mainly determined by the data and not the algorithm.

When data is the central role of the system, the quality of the data are very important. Poor data can lead to wrong assumptions, and have a negative effect on the application. Choosing the correct dataset is therefore a key factor [6]. Looking at the quality of the data can help you chose what dataset to use. Data quality can be described many ways, one of the more formal is from the ISO 8402 standard that describes quality as:

"The totality of characteristics of an entity that bear upon its ability of satisfy stated and implied needs" [7].

This indicate that data quality is something that is very depended on the intended application, and is therefore hard to generalize.

Quality of data is a subject that is gaining more and more attention due to the fact that the quantity of data available is larger than ever before. This forces researchers to choose between datasets, and a notion of quality in the dataset can help them choose. The heavy growth in available data is a result of projects that are moving data gathering from controlled labs, to the public. Many of these projects is citizen science projects, where it is the citizens who collect the data, and not the researchers [8]. This enables researchers to gather enormous amount of data, but they are no longer in control of the conditions the data is collected in, which introduces errors and other quality decreasing factors.

Non intrusive load monitoring is a topic that have been in focus in the past years. This is due to the rise of the smart meters, that makes it possible to measure at faster intervals, and collect the data on online services form real environments. But the smart meters architecture is designed after billing and regulation purposes, and not load monitoring. The network architecture is therefore often based on the unreliable user datagram protocol (UDP), since it is more important to get the current information fast, than get all information. This courses a lot of packets to be lost in transition, which can degrade the completeness quality of the signal. The missing data can be a problem for load monitoring, since many methods of load disaggregation is based on learning techniques. The quality of the signal can also help identify if the collected data is suitable as a training set.

2.1 Quality Criteria

It is not uncommon that different areas of research has its own quality criteria. This is due to the fact that quality is a very domain specific subject. One of the areas that have been dealing with citizen data for many years is the Geographic information area, that are used for maps, weather prediction and climate research. They have come up with several ways of describing quality in spatial data [9]. Method for defining quality in time series data have also been developed [10]. To better define quality criteria in the smart meter data inspiration from related work is used.

2.1.1 Related Work

Data quality is an area that recently have become a hot topic, due to the vast quantity of data. Many researchers strive to make tools that better can analyse data quality in different areas. In the area of spatial data is a "*Quality and Workflow tool*" being developed by the University of Wageningen [11]. The objective is to help researchers select the best suited data for a given data driven project. It does this by looking at different quality criteria, given by the user or found in standards for spatial data.

In bioinformatics is a tool named QCScreen developed to help create better dataset to metabolomics studies. In metabolomics studies is dataset often created by joining information from several different experiments of various quality. By using tools that can check the data quality and consistency to determine if a dataset is suitable for further processing, they are able to greatly improve the test results [12].

In the article *Taking a big Data approach to data quality in a citizen science project*[8] they talk about how quality assessment can be used to rate the believe on your data, and how to improve data collected in citizen science. The project focuses on bird observations, done by users on their smartphones. They improve the quality by disallowing the user to send incomplete datasets to the database, and in this way forcing the user to only deliver high quality information. They then cross check the information with information from people in the same area, to see if it varies greatly.

In the non-intrusive load monitoring (NILM) research area have data quality also emerge when talking about comparison between dataset. One of the tools working with this is the NILM-Toolkit (NILMTK) project [13]. In this they look at a series of quality parameter that can be used specifically for NILM. They point out that looking at the energy consumption observed by the main meter in relation to what is observed at the sub-meters tells a lot about the training set. This is since the success ratio of disaggregating appliances consuming lot of energy is higher than once that only consume a little.

One of the things all the methods have in common is trying to look at the completeness of the data. Some of the most low level criteria is the sample availability. The sample availability describes how many samples there is collected in relation to the expected collection amount, and look at how the samples are distributed in the measurement period. It is also seen that the activity is a good quality metric since a good data set must contain both areas with activity and areas without.

2.2 Quality In SmartHG Citizen Data

As a part of the SmartHG project 25 households have been equipped with meters on selected appliances and the main meter. The data collected from this experiment are prone with errors due to malfunctioning test equipment or unexpected interference from the resident which have resulted in offline measurement equipment for periods of time. Examples on unexpected interference could be if the resident is unplugging the measurable equipment, or turning off the power socket that supply's it. Unstable network does further degrade the signal, since the measurement equipment uses a lossy network.

2.2.1 Completeness And Activity

The SmartHG data is intended for appliance recognition, and the quality must be assessed with this in mind. The completeness and the activity in the data is therefore important.




Figure 2.1: 12 hour overview of house 10

On figure 2.1 is a 12 hour overview of the data in house 10 from the 16/8/2015. Here we see the sample availability and activity of the five different meters in the house. The availability is shown as a line where green is indicating that a sample is received as expected, and red shows a missing sample. On the figure it is shown that there is a few samples missing, which is to be expected due to the lossy network architecture. The Activity is also shown as a line, where green indicates activity and red indicates no activity. Activity is defined as a change in the signal, from prior values. From this we can see that the resident have a lot on activity on the TV from around 16:00 to 00:00, which we can presume means that the television is turned on in this period.

First the availability quality of the data is assessed. The availability quality for a specific period of time T_n for a specific meter m , is defined as the amount of samples observed in that timeslot over the expected sample amount. The resolution period T_P for each time period in T is chosen to be one hour.




Figure 2.2: Illustration of availability analysis

To calculate the amount of samples expected to be in a specific period of time $N_{max}^{(m, T_n)}$ does the sample phase $\phi_{start}^{(m, T_n)}$ for the given period need to be known. On figure 2.2 a illustration of a signal, the black dots are the samples and the red dots are the ones that are missing. The sample phase is the time from the beginning of the period T_n and to the first expected sample. This is needed since for a timeslot T_n of a length of T_P the maximum expected sample amount can vary with 1. This is shown in figure 2.2 where the period T_n has a potential of having 11 samples, where as period T_{n+1} only can have 10.

$$N_{max}^{(m, T_n)} = \lfloor \frac{(T_P - \phi_{start}^{(m, T_n)})}{T_s^{(m)}} \rfloor + 1 \quad (2.1)$$

$$q^{(m, T_n)} = \frac{N_{observed}^{(m, T_n)}}{N_{max}^{(m, T_n)}} \quad (2.2)$$

As shown on equation 2.1 is the maximum number of samples for a meter m in the period T_n calculated by taking the period time T_P , corrected with the sample phase $\phi_{start}^{(m, T_n)}$ for the given period, and dividing it with the sample time T_s . The quality of the meter is calculated as the ratio of observed samples in the timeslot T_n to the maximum samples, shown in equation 2.2.

To find the quality of a house in a given period T_n , that have a set of meters \mathbf{M} with a cardinality of M , we take the mean value of all the meter quality's, as shown in equitation 2.3.

$$\mu_{q(\mathbf{M}, T_n)} = \frac{1}{M} \sum_{m \in \mathbf{M}} q^{(m, T_n)} \quad (2.3)$$

A quality vector Q is constructed for each house. The quality vector contains the house quality

found with a period T_P on one hour. This have been done from March T_1 to October T_N .

$$Q^{(\mathbf{M})} = \{\mu_{q(M,T)} | T \in \{T_1, T_2, \dots, T_n, \dots, T_N\}\} \quad (2.4)$$

This is shown in equation 2.4 where \mathbf{M} is a set of meters in a given house. This can be graphically shown in figure 2.3 where all the houses Q vectors is shown. The color is a gradient running from light green for the best quality to red for bad quality.




Figure 2.3: Quality of houses in SmartHG project

On figure 2.3 is the availability quality shown for the 25 houses in the SmartHG project. When seen on this scale with a analysis resolution on one hour it is hard to see degradation caused by single sample missing here and there, but meters that have malfunctioned over a longer times shows itself. Since the quality of a house is the mean of several meter quality's, will this most likely appear as darker green spots, since not all meters are malfunctioning at the same time. But there still are some red areas indicating that all the meters in a house is not working.

There are the red spots that goes through all houses on the exact same time. This indicates that the server receiving the data for all the houses have been down, since it is unlikely that all meters in every house is down at the same time. The conclusion being that red dots most commonly are caused by the network being unavailable so the client can not sent to the server, or the server is unable to receive.

It is assumed that the first sample received from a meter happen at the time of meter installation, and the last sample received is the time of meter removal. The meter is assumed to be operating in between these two points in time. On the figure does the coloured Q vector starts at installation

time, and ends at removal time. This illustrates how some houses have been operational longer than others.

Since the data is intended for appliance recognition it is of interest where in the data there is activity, and where there is nothing happening. Both areas are important for a NILM application in training scenarios. We define activity as area in the data where there is change as described in equation 2.5.

$$f(x) + \epsilon < f(x+1) \vee f(x) - \epsilon > f(x+1) \quad (2.5)$$

Where ϵ describes a threshold to filter out changes caused by noise. This can also be described as the standard deviation over an area is greater than the threshold. The activity is analysed in the available data, and is shown on figure 2.4 where green is high activity and red is non activity.




Figure 2.4: Activity of houses in SmartHG project

The activity shown on figure 2.4 is the average of the activity of the meters in the house. There is almost never a meter that does not have at least a little activity in a house, so a complete read area is fairly rare. The most interesting areas in the activity map, is often the places where there is a lot of change in the amount of activity like for house 16.

2.2.2 Main Meter In Relation To Sub-meter consumption

As discussed in the NILMTK project [13] is the relation between the energy consumption measured by the main meter and the consumption measured by the sub-meters an important quality metric. This can give two important indications about the dataset. The first indicator is how

much knowledge about a house we have. A house where only a small fraction of the energy consumption is accounted for can give problems when applying NILM algorithms, since there can be a lot of unknown noise from other appliances. The other is more appliance specific. If a appliance is not responsible for a significant part of the total consumption, odds are that it is almost never on, or that it usages so little energy that it is hard to detect. Both can have a significant impact on NILM algorithms. This will be further discussed in chapter 4 and 5.




Figure 2.5: Usage of SmartHG houses

On figure 2.5 is the total consumption of the houses shown, and split up in the the different sub-meters, and a "other" category. The "other" category is the energy the main meter has used besides the energy accounted for by the sub-meters. On the figure is only shown the houses which main meter had observed a energy equal or greater than the sum of all the sub meters. This is not the case for houses that have malfunctioning main meters for a longer period of time, or for houses that have generated a lot of energy form solar cells.

As seen on figure 2.5 is the appliances equipped with sub-meters only responsible for a minor part of the total house consumption. Only the TV's and some of the unknown devices, have a relative high energy usage. The unknown devices is devices that have a sub-meter, but the appliance on the sub-meter is unknown.

2.3 Chapter Discussion

When working and gathering big amounts of data it is important to be able to gesture about the data prior to analysing it. To look at different quality measurements researchers are able to get a good general overview of the data, and find better areas to focus on.

It is shown that the quality also can indicate problem in the data gathering phase of a project. problems such as malfunctioning equipment, network failure or server problems can be detected by quality monitoring systems.

For NILM applications the activity in the data, and error rate in the data can be good indications of quality. Further more is the amount of energy observed by the main meter in contrast to the energy observed by the sub-meters in a house also a good indication of the dataset quality.

Gap Reconstruction 3

One of the more common problems in citizen science projects is gaps in data. This can happen either if the network connection is unstable or the test equipment gets prematurely turn off, as discussed in chapter 2. This can greatly degrade the data quality, and lead to errors in the application. One way to deal with this problem is to use mathematical gap filling techniques to come with a qualified guess on how the data would look like in the gap.

In order to use this methods we must assume that the missing data in the gap follows the same behaviour as the data on each side off the gap. If the signal is so stochastic that this is not the case then gap filling is not recommended[14].

In the case of the SmartHG project the data can be seen to have a part that is depended on the previous and future data plus a stochastic part that are determined by the user and the appliance. Due to the stochastic part a perfect reconstruction is not possible, but it is the hypotheses that the non stochastic part is still so dominant that a decent reconstruction is possible.

3.1 Gap Filling Methods

Various methods exists for gap filling. Five popular algorithms are selected for this project, and is validated on the SmartHG project data. These methods have been chosen since they all have a different approach on the gap filling process. For some of the algorithms it is important to keep the frequency spectra as intact as possible, for other it is the jitter power or the exact sample value they are trying to estimate. Some algorithms have also been designed for small gaps, while others have been designed for larger.

In the following sections it will be briefly described what makes the five algorithms unique, and how they work.

3.1.1 Papoulis-Gerschberg Algorithm

The Papoulis-Gerschberg (P-G) algorithm is a multi gap filling algorithm, meaning it is capable of correcting more than one gap at the time. This makes the algorithm preform good in conditions with many gaps and few available data points between gaps. This is due to its ability to collect information about the signal across multiple gaps[15]. The P-G algorithm works under the assumption that the signal is a periodic stationary signal with a known bandwidth. The signal will therefore consist of M frequency components, and everything outside the band is assumed to be noise. The signals in the SmartHG is not stationary, but for small snippets can approximately stationariness be assumed.

The true bandwidth is also unknown in the signal. The P-G algorithm is very depended on the

bandwidth for a correct reconstruction. A modified version of the algorithm that estimates the bandwidth, by varying the frequency components M and analysing the mean square error on the known signal is therefore used [16]. This approach is fairly good at estimating the true value of M , but it is time-consuming.

3.1.2 Wiener Filling Algorithm

The Wiener filling algorithm is an extension of a Wiener predictor, the Wiener predictor assumes that there exist a linear relationship between the next sample and the previous samples. By trying to predict the missing samples from both sides of the gap, and combining the knowledge, it estimates the missing samples [17]. For larger gaps does this methods rely on earlier predictions to close the gaps. This result in errors being accumulated over the gaps. The method is fast, and is therefore suited for large data with small gaps.

3.1.3 Spatio-Temporal Filling Algorithm

The Spatio-Temporal filling (SSA) algorithm uses singular spectrum analysis to split the signal into a series of sub-signals. The sum of the sub-signals is the original signal, and the sub-signals are ordered so the most dominant is first, and the least dominant is last.

The reconstruction philosophy is that the gap has introduced noise in the signal, but a sum of only the most dominant sub-signals must be close to the original signal without noise. But in order to know how many sub-signals to include in this sum, we introduce an other artificial gap. While the sub-signals are being accumulated the mean square error of the artificial gap is observed. When this mean square error hits its minimum peek, it is assumed that the reconstruction is as good as possible [18].

This method is very popular for gap filling. It has shown to be very noise resistant since it finds the overall trends in the data. It does require quite a lot of data to be known post and prior to the gap since an artificial gap must be introduced. It is based on singular spectrum analysis which assumes that the signal consist of stationary processes. This is a similar constraint to the P-G Algorithm in section 3.1.1.

3.1.4 Envelope Filling Algorithm

Unlike the previous described methods does the Envelope filling algorithm not depend on frequency analyses, but rather on the expected power of the signal. Looking at the envelope of the signal it assumes that all local maxima and minima must lie on the upper and lower envelope. It then looks at the data prior and post the gap and try to estimate the number of local maxima and minima in the gap, and their locations. It does this by looking for patterns in the time series data [10]. When the new maxima and minima are found the points is connected by using spline [19].

The methods do not make any assumptions about the signals stationariness or bandwidth. The method can also be used on none equally spaced time series.

3.1.5 Empirical Mode Decomposition Filling Algorithm

The empirical mode decomposition filling (EMD) algorithm uses empirical mode decomposition, to break the signal into intrinsic mode functions (IMF). The sum of all IMF's is the original signal. The IMF's is all more low frequent and simpler in structure than the original signal. The hypothesis is that it is easier fixing a gap in a simple signal than a complex one.

The envelope filling algorithm in section 3.1.4 is used to fix the gaps in the IMF's. The IMF's can now be accumulated to get the original fixed signal. Like the envelope filling algorithm does it not make any assumptions about the signals stationariness, bandwidth and can be used on none equally spaced time series. But making a empirical mode decomposition on a signal with a gap in is a non trivial process and can introduce errors [19].

3.2 Gaps In SmartHG Dataset

The gaps in the SmartHG project dataset is caused by a lot of different sources such as bad network connection, unplugged measurement equipment or server breakdown. This makes the type of gaps different from case to case. Three aspects of a gap is important for the gap filling: The size of the gap, the data known before the gap, and the data known after the gap.

3.2.1 Gap Size

On figure 3.1 is the quantity of different gaps shown. The different gap size is measured as samples missing, e.g. a gap size on 3 means that there is 3 samples missing in a row, between two received samples. Looking at the different gaps in the dataset we see that the normal gap is relatively small. Most of the gaps are between 1-5 samples as seen figure 3.1.




Figure 3.1: Gap quantity




Figure 3.2: Error recovery capability

The aim of data reconstruction is to recover the missing samples. The gap size correction capability is a metric telling how big gaps it is possible to correct. If a application has a gap size correction capability on 5 it means that is is capable of correcting gaps that have the gap size of 5 samples or smaller.

On figure 3.2 is shown how much of the missing signal that can be reconstructed with different gap size correction capability. It is shown that a if a gap size correction capability of approximately 20 samples can be achieved, it is possible to recover 30% of the missing data in the SmartHG

dataset. Since the signal is partly stochastic, and it is not possible to recover the stochastic part, the complete signal can never be reconstructed. The greater the gap, the greater influence does the stochastic part have on the signal. Smaller gaps can therefore be fixed with greater success. It is therefore unlikely that recovery of more than 30% will be possible.

3.2.2 Post And Prior Knowledge

The reconstruction process works by looking at the samples available prior and post of the gap. Common for all reconstruction methods is that they assume that the signal in the gap must have behaved in relation to the samples prior and post for the gap. Therefore is the samples prior and post for the gap called knowledges, since it grants the knowledges used for reconstructing the gap.

But in a signal with lots of gaps it can be interesting to see how much knowledge is available for reconstructing a gap. This is done by seeing how many samples that are available prior and post to the gap. This is important since much data allows for detailed models, that can improve gap reconstruction and little knowledge gives the stochastic part dominance which will lead to error prone reconstruction.




Figure 3.3: Available samples for recovery

In the case of the SmartHG project dataset the data available prior and post to a gap varies greatly but is around the same max and min values for every gap size. The median samples available to fix a gap is around 6 samples as shown in figure 3.3. On the figure is shown that no matter the gap size is the expected knowledge about the same. This can be problematic since larger gaps offend needs more knowledge than smaller gaps. This further indicates that complete recovery of all gaps is impossible.

3.3 SmartHG Dataset Reconstruction

The data reconstruction algorithms mentioned on section 3.1 have been tied on the SmartHG dataset. First have 6 unique error free areas in the data been found, and a artificial gap have been introduced in the 6 scenarios. The gaps have now been reconstructed, and a comparison to the true value is made.

The 6 scenarios have been randomly chosen under the constraints that there where no missing data in them, and they all are different in activity level from the other chosen scenarios. Both accumulative scenarios and non accumulative scenarios have been chosen.

There are several ways of comparing the different algorithms, in this report 3 have been selected based on the likelihood of the importance in a learning algorithm, which is the target application for the data. The first is a simple sample by sample comparison, where it is seen how much each sample differs from the true value. The other is the frequency, where it is analysed how much the frequency response is different. The last is a method called jitter compression, where we see that the power of the jitter is the same, more on this in section 3.3.3.

3.3.1 Sample Comparison

The sample by sample comparison is created by finding the mean square error of the samples compared to there true values.




Figure 3.4: Sample comparison of the reconstruction methods

On figure 3.4 is the result shown for 3 different cases. The Known attribute indicates how much prior and post knowledge in samples is used to reconstruct the gap. On the x axis is shown the gap size, as expected does the prediction grow worse the greater the gap size. As a baseline is a linear predictor also added. The linear predictor makes a simple line between the two samples at each edge of the gap, and places all missing samples on the line.

As seen on the figure the best reconstruction is made by the P-G algorithm, with relatively few samples as knowledge. The wiener algorithm is also close to the linear base line. The logical statement is the more information or knowledge we have the better should the prediction become. This does not seems to be the case, since a lot of the algorithms amperes to preform better at lower knowledge rate. This can be explained by the stochastic part of a signal, when the knowledge is small it is more likely that the information around the gap is relevant, the more knowledge grows the more likely is it the a stochastic event will bring false information to the model. This can be illustrated as on figure 3.5.




Figure 3.5: Stochastic effect on prediction.

On the figure 3.5 is shown a original signal with a gap, indicated by the red area. The true value is illustrated by the full red line in the original signal. The two subsequent graphs is the reconstruction, the blue area is the known area and the dotted red line is the predicted values. When the known area is small there is not enough frequency information to make any grant changes, so the prediction will look more or less like a linear interpolation. This is compared to the true value not a bad guess. When there is more samples known we are able to make more complex estimates of the signal. On the figure it is shown how when we have a knowledge of 15 we also see the stochastic event. Taking this in to the model, changes the prediction to a less correct value, due to the frequency's added by the stochastic event.

Since the results shown in figure 3.4 at higher knowledge shows that the reconstruction methods for the most part preforms worse than the linear basecase, we can conclude that the signal is quite stochastic and not as periodic as one would hope. For many smart-meters it is possible to get both the instantaneous current power and accumulated power over time as the samples. By looking at the accumulated power before and after a gap it is possible to gesture about the changes in power during the gap. This information can be used to further improve the gap filling.

3.3.2 Frequency Comparison

An other metric to compare is how well the frequency response is preserved in the reconstruction, since this is often more important than the true value itself for smaller devices with many states like computers. Such devices are often recognised by their usage pattern and not the true power draw. By looking at the Fourier transform of the original signal and the reconstructed power of the different frequencies was compared.




Figure 3.6: Frequency comparison of the reconstruction methods

As shown on figure 3.6 both the Wiener and the P-G Algorithm seems to perform well. Like for the sample comparison the most successful reconstruction seems to be at low knowledge.

3.3.3 Jitter Comparison

The jitter is a different metric as it tells how well we have modelled the "noise" of the signal. In a jitter analysis the signal is thought of as a slowly changing signal with a faster noise signal embedded in it.

The slow signal is assumed to be found by taking a low pass filter over the signal to filter out the noise. The jitter analysis measures the power in the jitter from this signal as illustrated on figure 3.7.




Figure 3.7: Jitter Power Illustration

It the estimated jitter power have been compared with the real jitter power for all the reconstruction methods.




Figure 3.8: Jitter Power comparison of the reconstruction methods

As seen on figure 3.8 does non of the methods model the noise very well, as is to be expected due to the nature of the selected methods. This is will probably not be a problem for the SmartHG data, since it is sampled at a very low sample rate, and noise therefore is not likely to be used as a parameter. For data there have been sampled with high frequency it have shown that the noise created by the different appliances can be used to detect them [1].

3.4 Chapter Discussion

In the chapter the errors of the smartHG dataset is analysed. The dataset is a raw dataset, that have not undergone any cleaning prior to the process. It is shown that the most common errors in this types of data is single sample errors, and minor gaps of up to 5 missing samples.

different kind of gap filling techniques is used on a series of artificial introduced gaps. It is shown that the small gaps, that are the most common, can be fixed relatively easy. For larger gaps it is harder to do, due to the stochastic nature of the signal. The stochastic element of a electric signal is often greater than the periodic part, which is a challenge for the gap filling. This combined with the low sample rate makes it hard to gesture about the true frequency content of the signal, since it is sampled way below the Nyquist rate.

Appliance Recognition 4

Appliance recognition and load disaggregation are some of the key aspects in non-intrusive load monitoring (NILM). The main purpose of the NILM approach is to better understand the power usage in the home. This information is used by the electric companies to better plan the energy production and distribution. It also have the potential to help the different households to more optimal power savings. This is done by informing the residents about the power consumption of individual appliances. It is shown that informing a household about its usages pattern can lead to significant savings [5]. This makes the basis for load disaggregation, where appliance specific load is guessed based on the main meter load. This enables the user to know what uses the energy and when.

One other aspect of NILM is event detection. Sometime the actual consumption of an appliance is not as important as when it was used, and for how long [find a ref](#). This is a simpler task than guessing the true consumption, and for some applications is a better approach. As an example is this kind of information sufficient if we want to track the activity of an elderly person in an assisted living situation.

4.1 NILM Concepts And Challenges

The aim of NILM is to partition the whole house consumption data in to appliance specific consumption.

$$P(t) = p_1(t) + p_2(t) + \dots + p_n(t) \quad (4.1)$$

This can be seen as in equation 4.1 where $P(t)$ is the total consumption at time t , and $p_n(t)$ is the consumption of appliance n at time t . The aim is to partition $P(t)$ back to the different appliance specific consumptions. In order to structure the problem we categorise the appliances in 4 groups[1]

Type-I: Appliances that only have 2 states corresponding to on/off. This could be a lamp or a water boiler.

Type-II: These are appliances with multiple states, and can be modelled as a finite state machine. Many modern devices such as TV's, computers and washing machines fall in this category.

Type-III: These appliances are referred to as "Continuously Variable Devices". These devices have a variable power draw and is impossible to model as a finite state machine. This could be appliances like power drills, and dimmer lights. These are by far the hardest for the NILM algorithms to detect.

Type-IV: These are a special kind of appliances that are always on and consume energy at a constant rate. Such devices could be smoke detectors and TV receivers.

Some of the different appliance types is illustrated on figure 4.1.




Figure 4.1: Appliance types. Source [1] [make your own figure](#)

The appliances in Type-IV is not a particular researched area, since it does not give the user much information to calculate savings from. These appliances are often using very little energy, and must not be turned off at will.

Type-I and Type-II are very used, and can cover almost any appliance. In the area of event detection is the most common approach to see all appliances as Type-I and detect on/off events. Most appliances today is Type-II due to the growing amount of complicated electronics that are embedded in them. Then there are devices that are Type-II but most of the time act like Type-I appliances. An example of this is the vacuum cleaner. On most vacuum cleaners you are able to change the suction intensity, which makes it Type-II. But most people do not change this very often, so the energy usage looks more like a Type-I appliance.

4.1.1 NILM Features

For data disaggregation it is common to use method based on machine learning and optimization. The approach is to first extract features from the dataset, and then train or validate advanced statistical models by using these features.

In NILM the features can be sub categorised in steady state, transient state and non-traditional features. These categories can be further expanded as shown in figure 4.2.




Figure 4.2: Feature types. Source [1]

Steady state features are features that can be extracted when the signal is in the stable state. Features that are often extracted in this state are "Power Change" which is the jump in power or reactive power usage. Time and frequency analysis is often used on the voltage and power signals of the meters. It is shown that analysing the harmonics in signals is a powerful way of making appliance detection [20]. The trajectory between voltage and power usage is a good feature. This can be used to detect the reactive power of a signal, which is a give-away for inductive loads. Some researchers have shown that different appliances make different noise profiles on the main line. Unfortunately, does this kind of recognition require an extremely high sampling rate [21].

The transient state of an appliance is a short state that comes when an appliance switches between steady states. In this state there often are appliance semi-unique waveforms in the current and voltage domain. The noise generated on the mainline in this state is also a good indication of the appliance type.

There also exists a series of non-traditional features, that is used in special types of algorithms. One of them is based on matching the power usage profile to geometrical figures, and using the series of figures to recognize an appliance [1].

Many of the above features require a high sample rate in kHz or MHz span in order to be effective. This kind of resolution is often not available since it is normal smart meters that are providing the data. For a smart meter a sample rate at 1 Hz would be considered fast, and for many applications this is down to 0.03 Hz or slower. It is therefore common only to use the steady state features, and most of the time only the power change feature, which looks at changes in power, reactive power, current or voltage.

4.1.2 Learning Strategy

When you have extracted features from the signal there are several algorithms that can be selected for the training process. They can roughly be split up in two categories: optimization algorithms and pattern recognition algorithms. The optimization algorithms rely on a database

of appliance info, and tries to find the subset of the database that are most likely to be in use.

$$M = \underset{x \in \mathbb{P}(D)}{\operatorname{argmin}} \left(\left| \sum_{i=0}^I x_i - \hat{y} \right| \right) \quad (4.2)$$

This is illustrated in equation 4.2 were a set of appliances M that correspond to the measured signal \hat{y} needed to be found. This is achieved by searching the known database of all appliances D . This is done by taking the powerset of the database and finding the subset that minimises the error. It can therefore be seen as a minimization problem, where the desired set M is the arguments to the minimum solution. This approach is fine for a small amounts of appliances, but for large databases and many appliances in use at the same time this would take far too long to be practical.

Due to this, most researchers use pattern recognition techniques instead. These have the benefits of being more scalable, and be usable even if only parts of the device database is known. There are generality two approaches to pattern recognition: supervised and non-supervised. In supervised training it is assumed that the ground truth of your system is known. For a NILM application this means that the true consumption of the appliance need to be available for the training process. This means that a sub meter must be placed to measure the appliances individually, which often is costly and time consuming. In non-supervised methods only the main meter readings are available, and no information about the appliances are known. This type of algorithms will usually try to cluster the readings in distinct groups cosponsoring to a guess of what appliances that correspond to the reading. Fore this kind of training it is easy to collect a big training set, but it requires the groups to be manually labelled afterwards.

Most algorithms are based on clustering techniques and hidden Markov models (HMM) [find a ref](#). HMM have proven very effective when classifying Type-I and Type-II appliances. This is due to its ability to model state changes as an imported part of the classification. It is shown that using clustering techniques can greatly improve the classification, and reduce the training time. Artificial neural networks are also showing great promise, and is a topic that is currently in focus by many NILM researchers. More on this can be found in section 4.2.

4.1.3 NILM Challenges

There are many challenges in NILM. One is the Type-III appliances. Due to their varying, and at times random, energy consumption it is hard to fit them to a behavioural model. This is made further harder by the low sample rate. Many great features of the appliances is hidden due to the low sample rate most data is collected with[1].

A other problem is the "heavy consumer problem". In a household there are some appliances that require a lot of energy when they are on, such as stoves, air conditioners and refrigerators. These are often refereed to as the heavy consumers. Then there is the appliances that consumes very little energy like laptops, DVD players and routers. The heavy consumer problem is that the heavy consumers is so dominant that they make it hard to track the power signature of little

consumers. This is why many researchers only focus on the top 10 heavy consumers when during NILM applications [13].

4.2 Related Work

There are several interesting problems and approaches being researched in the NILM community in the moment. One is the lag of a good bases of comparison, both to compare the performance and the training time of a solution. It was a common practise that each researcher uses his own dataset, or an artificial lab setting to validate the results. This made it hard to compare the algorithm to others, and raised questions about the correctness of the findings.

To accommodate this problem the NILM-Toolkit (NILMTK) was created. This is a collection of python libraries, that create a framework to evaluate NILM solutions. It supports a wide range of known datasets that can be used to train or validate an algorithm. It also comes with two benchmark algorithms to be used to compare a new algorithm against. The NILMTK will help researchers to easily share and compare algorithms in the future[13].

Another framework to accomplish better comparison and sharing is the NILM-eval framework. This is based on Matlab and supports algorithms to be written in either Matlab or in python. The NILM-eval and the NILMTK is created to solve the same problem, and is compatible with each other in the sense that the interface for algorithms are the same[4].

Deep neural networks have last years improved the area of computer vision and speech recognition remarkably. NILM researchers are starting to look into if similar techniques can be used to improve data disaggregation. Neural networks have like HMM the ability to model state changes. Furthermore they are capable of automatically finding important features in the data. Preliminary tests shows that a deep neural networks can outperform the classic HMM approaches. The downside with the deep neural network is that it takes weeks to train, and there are thousands of parameters there can be changed. This makes the task of finding the global minimum hard, and the method is very prone to be stuck in a local minimum. The method is currently being used on heavy consumers[22]. Current research is looking that its potential in solving the heavy consumer problem, but no definite conclusions is yet to be reached.

Many approaches is supervised, since this kind of approach is better at separating the appliances. But this require that you first build a database with information about all the appliances. To have this prior knowledge of the system seems unrealistic in an application setting, where it is most likely you do not know anything about the house. To combat this, various unsupervised methods have been developed. The focus in this kind of studies have been load separation for the sake of making power savings, or to inform the electric companies to help them create a more stable electric grid. There are several good algorithms for this today that support this purpose. For event detection and user habits analysis it is still most common to use the supervised approaches [23]. This report will therefore focus on the supervised approaches.

4.3 Recognition Methods

Methods based on HMM are currently the most popular method for NILM. Three method is selected to be validated in this paper. The methods are selected since they are often referenced in the literature, and is often used as benchmark algorithms. All of the algorithms are based on HMM or clustering in different configurations. The following subsections will briefly introduce the algorithms.

4.3.1 Factorial Hidden Markov Models

The HMM has proven to be one of the most used tools to create probabilistic models based on time series data. This is done by seeing the system as a series of observable states Y_t , and a series of hidden states S_t [2]. It is assumed that there exists a probabilistic relation between the observations Y_t and a sequence of hidden states S_t . This is implemented as a first order model, which implies that the current hidden state S_t is only depended on the previously state S_{t-1} and the current observed state Y_t .

$$P(\{S_t, Y_t\}) = P(S_1)P(Y_1|S_1) \prod_{t=2}^T P(S_t|S_{t-1})P(Y_t|S_t) \quad (4.3)$$

The joint probability for a specific hidden state S_t that has caused the observed output Y_t can be seen in equation 4.3. $P(Y_t|S_t)$ is the emission probability, which shows the conditional probability for Y_t being outputted by a given hidden state S_t . The probability for a state change is given by the transition property $P(S_t|S_{t-1})$.

Factorial hidden Markov models (FHMM) is an extension of this classic HMM where there exists several hidden state chains that all contribute to the observed output. This can be seen as an extension to equation 4.3 where the hidden states S_t can be seen as a set of states as shown on equation 4.4 where M is the number of chains.

$$S_t = \{S_t^{(1)}, ..., S_t^{(m)}, ..., S_t^{(M)}\} \quad (4.4)$$

This can also be shown graphically on figure 4.3. It is illustrated how several hidden states from different chains contribute to the output.




Figure 4.3: Factorial hidden Markov model illustration. Source [2]

As an extra constrain can each hidden state, in each chain, only depended on the previously state in the chain, and not on the hidden state of other chains. This makes the state transition probability to be defined as in equation 4.5.

$$P(S_t|S_{t-1}) = \prod_{m=1}^M P\left(S_t^{(m)}|S_{t-1}^{(m)}\right) \quad (4.5)$$

The observed stated is therefore created by the contribution of many different hidden states.

In this implementation there is normally used only two or three states, since most appliances are seen as Type-I or Type-II appliances. One chain for each of the appliances plus one that is called the other chain is created in the model. The other chain is to allow for the other appliances that are not taken in to the model and noise. The model can therefore be used to validate for a given output what is the probability that a given appliance is contributing to the output. The concrete implementation for this algorithm has been borrowed from the NILMTK [13].

4.3.2 Parson

The Parson algorithm is a variant of the Kolter and Jaakkola algorithm[24]. The algorithm is designed to be a semi-supervised learning algorithms. This means that the algorithm does not need to train on the data from the concrete house prior to deployment, but it still requires some prior knowledge. The idea is to have a database of general appliance models, that tells how specific appliances most likely will behave. This model will be able to detect the device, and do some unsupervised special training to better fit the concrete appliance.

The algorithm is based on a kind of HMM called "difference HMM" since they take the step difference in the aggregated data and use as the observed output Y_t . In a normal HMM there is

only emitted one output for each hidden state in a chain. In the Parson algorithm the model is changed to output two states from the hidden states X_t and Y . This creates a model structure that can be graphically illustrated as in figure 4.4[3].




Figure 4.4: Parson hidden Markov model structure. Source [3]

The Y_t is corresponding to the step in power that is observed since the last sample. It is assumed that the only thing that can cause a step in power is if an appliance changes states from one to another. The X_t is the constant power consumption the appliance would have in the state. This is used to filter out the appliance if the power consumption observed is smaller than this threshold.

It is assumed that the power consumption of a state in the general model of an appliance can be modelled as a Gaussian. This can be illustrated as in equation 4.6.

$$w_t|S_t \rightarrow \mathcal{N}(\mu_{S_t}, \sigma_{S_t}^2) \quad (4.6)$$

This shows how the power consumption of an appliance is modelled as a normal distribution. The step Power Y_t , can also be modelled as a normal distribution that is created from the change between two states as shown in equation 4.7.

$$Y_t|S_t, S_{t-1} \rightarrow \mathcal{N}(\mu_{S_t} - \mu_{S_{t-1}}, \sigma_{S_t}^2 + \sigma_{S_{t-1}}^2) \quad (4.7)$$

To model the constraint that the meter only can be on if the power consumption present in the signal is higher than the requirement for the state is an additional constraint added to the

emission probability. This constrain is mathematically described in equation 4.8.

$$P(w_{S_t} \leq x_t | S_t) = \int_{-\infty}^{x_t} \mathcal{N}(\mu_{S_t}, \sigma_{S_t}^2) dw \quad (4.8)$$

Where x_t is the measured power consumption and w_{S_t} is the appliance power consumption constraint. This makes the probability going towards zero if the total power draw is too little and towards 1 if there is more than enough power. This is a loose constraint since too high power draw also could be created by other appliances. To minimise the impact this have for the algorithm is the mean power draw subtracted from the data before searching for other appliances.

In the implementation validated in this report are all appliances modelled as Type-I. The algorithm is allowed to train on the house data, to better fit the general models to the appliances and create the initial database.

4.3.3 Weiss

The algorithm proposed by Weiss called "AppliSense" is a other approach that do not use Markov models[25]. The philosophy is that some appliances like lamps and kettles are purely resistive, where others like washing machines and air conditions are more inductive and some appliances like laptops are more capacitive. By measuring not only the power consumption, but also the amount of reactive power it is revealed if there is inductive or capacitive appliances.

The algorithm is supervised since it requires a signature database to be created prior to usage. The signature database consists of how the changes will be in reactive and real power for a specific device when it is turned on and off. In their paper the authors show a method where this database can be created of a user, by using a smart-phone application and turning on and off appliances in the house, and inputting the information in the mobile application[25]. In this paper the database is created by data collected from the sub-meters measuring only the devices.

The algorithm looks for significant power changes in the data. If a change is found it is assumed to be an on or off event. The change is now compared to all the values in the database. If a close match is found it is assumed to be this device, if not it is a unknown device. This limits the approach to Type-I appliances, and the algorithm requires the reactive power to be measured.

4.4 The ECO Dataset

The experiments in this chapter will be done on the Electricity Consumption and Occupancy (ECO) dataset [4][26]. The dataset consists of electricity consumption measured from six households that have been equipped with sub-meters on selected meters and the main meter. The data is collected in Switzerland in the period June 2012 to January 2012.

The data is sampled with a resolution on 1 Hz on both the sub-meters and the main meter. The main meter data contains information about real and reactive power, where on the sub-meter only real power is measured.




Figure 4.5: House energy distribution. Source [4]

On figure 4.5 is shown how different appliances are responsible for the energy consumption in the different houses in the ECO dataset. For house two, more than 75% of the energy usages is accounted for by the sub-meters. This makes house two a preferred candidate for experiments, since we know who the heavy consumers are and who the little consumers are.

4.5 Validation of Methods

This section describes the experiments conducted using the three methods mentioned in section 4.3. The experiments is validated using F1 and accuracy sensitivity and specificity metrics's. Since the focus is event detection a sample can fall in one of four categories.

True Positive This is when the guess says that there is an event, and there really is an event. Corresponding to the appliance is on.

True Negative This is when the guess says that there is no event, and there really is no event. Corresponding to when an appliance is turned off.

False Positive This is when the guess says that there is an event, but there really is no event.

False Negative This is when the guess says that there is no event, but there really is an event.

The F1 and accuracy score is calculated as shown in equation 4.15 and 4.16

$$TP = \text{number of true positives} \quad (4.9)$$

$$TN = \text{number of true negatives} \quad (4.10)$$

$$FP = \text{number of false positives} \quad (4.11)$$

$$FN = \text{number of false negatives} \quad (4.12)$$

$$recall = \frac{TP}{TP + FN} \quad (4.13)$$

$$precision = \frac{TP}{TP + FP} \quad (4.14)$$

$$F1 = 2 \times \frac{precision \times recall}{precision + recall} \quad (4.15)$$

$$accuracy = \frac{TP + TN}{TN + TP + FN + FP} \quad (4.16)$$

The F1 score is mainly focused on how well the system guesses the the events, and do not score the true negatives. The accuracy is slightly different in that it looks at how many correct guesses there are in relation to the total amount of guesses. For a NILM application the accuracy score on its own is not a very representative metric since some appliances have way more off time than on time [27]. To guess that a appliance is off is just as good at guessing that it is on for the accuracy score. Appliances that are mostly off can therefore get a high accuracy score even though it is hard to determine when they are on. Therefore is it necessary also to look at the F1-score. This score tells how accurate it can be guessed when a appliance is on. This is done by combining the recall, that tells how many true positive was guess correctly in relation to the total amount of positives, as described by equation 4.13, and the precision that is the ratio of correct guessed positives events in relation to all positive guesses, as shown on equation 4.14.

Five appliances have been selected from the ECO dataset house 2. The appliances have been selected so there is some heavy and some little consumers. The appliances selected are: water kettle, fridge, TV, stereo and Laptop. All the methods will be traind using a training period of 15 days and validated on 75 days of data.

4.5.1 Sample rate experiment

Common for the three methods is that they are designed for signals sampled at low sample rates in the sub 1 Hz rate. To determined the sensitivity, on the sample rate, is an experiment conducted where the dataset is downsampled to simulate different sampling rates. The downsampling is done by applying a mean filter, so the new sample will be the average values of the samples in a specific period. This method is selected since it is the assumed behaviour of a meter that transmits samples at low rates.

The results shown in figure 4.6 illustrates the average score of the different algorithms when recognising the five appliances discussed in section 4.5.




Figure 4.6: Algorithm validation at different sample graduality.

Both the accuracy and F1 score is shown. It is shown that for the Weiss clustering algorithms the overall performance is decreasing for both scores. This indicates that it is both harder for it to determine when an appliance is on and when it is off. There is all most no change in the F1 score for the Parson and FHMM algorithms that both builds on HMM. Where the accuracy is decreasing for the parson as well. This tells that it finds the areas where the appliance is on just as well at high sample rate as for low. But to correctly guess the off periods is getting harder for the parson algorithm, where the FHMM algorithm seems to be unaffected.

It seems that for high sample rates does the Weiss algorithm preform superior to the others, where for lower is the Parson algorithm that preforms well. The FHMM algorithm seems to preform worse that the others, but is more robust to changes in the sample rate.

The results shown in figure 4.6 is the average for all appliances. The different appliances individual scores shows that some appliances are easier to find than others. The accuracy score for the different appliances is shown in figure 4.7.




Figure 4.7: Appliances accuracy score

Since the accuracy score tells how well all events are guessed it can be heavily dominated by the off events. Most of the appliances selected is mostly off appliances. Especially the water kettle is a good example as it is only used a few minutes each morning, and is off rest of the time. If it was assumed that this device was always off it would get a fairly high score since it is off most of the time.

This is the reason why the score is fairly high for all the different appliances. It is also shown that the top algorithm is different from appliance to appliance. The fridge and the stereo is the only appliances that are greatly affected by the change in sample rate, the rest seems to be relatively stable.

Since the accuracy score is dominated by the off periods, the F1 score is also shown in figure 4.8. This once again indicates how the different appliances has different algorithms that best detect them.

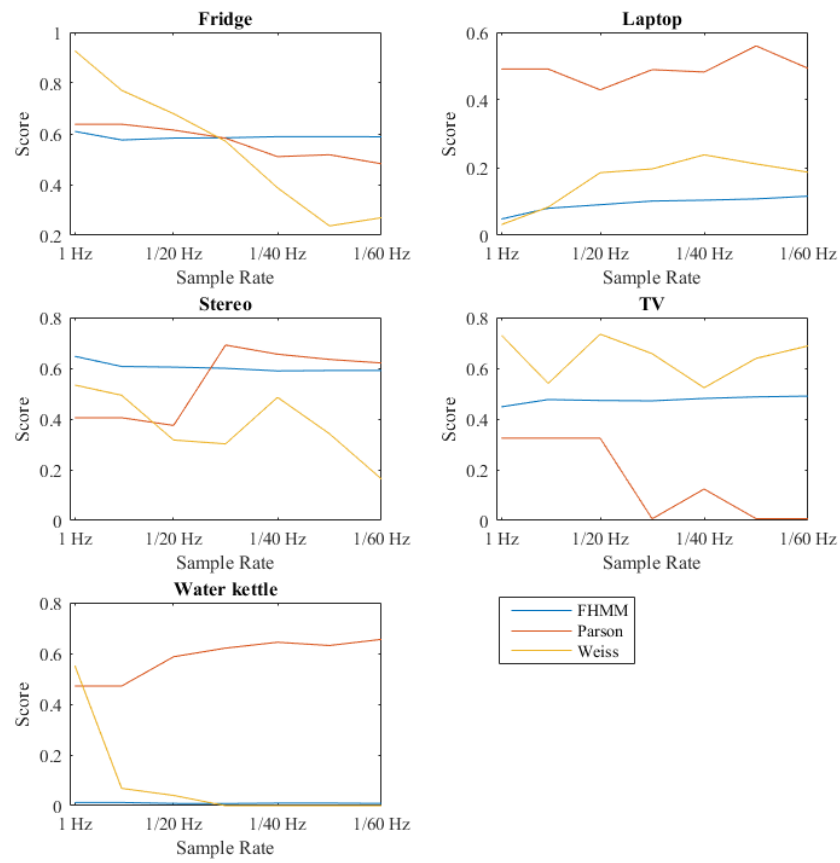


Figure 4.8: Appliances F1 score

4.5.2 Error Tolerance Experiment

As discussed in chapter 2 is the data collected often filled with errors. This degrades the quality of the signal which can lead to worse performance. To validate the errors impact on the algorithms is an experiment conducted where the error rate is increased, and the F1 and accuracy score is measured.




Figure 4.9: Algorithm validation at different error rate

The experiment was preformed by randomly introducing errors with a probability given by the error rate. The errors was introduced only on the validation dataset, to simulate a situation where a model was created in an error free environment, and then deployed in an error full environment. The errors is introduced independently of each other, simulating network packet loss error. This is the most common error type in the SmartHG dataset, as discussed in chapter 2. This was done up to 30% where approximately $\frac{1}{3}$ of the samples are missing. The results is shown in figure 4.9. It is shown that all methods is impacted by the errors, and have a decrease in performance. The experiment is conducted at a simulated sample rate on $\frac{1}{30} Hz$.

The algorithms tries as default to make a mean filtering to try to remove errors and noise. This will in situations with many errors remove some of the steps from the signal, and decrease the performance.

4.5.3 Gap Filling Experiment

In chapter 3 it was described how mathematical gap filling methods could be used to estimate values for the missing gaps. It was mentioned that due to the stochastic nature of the signal, it was impossible to make a perfect reconstruction, but a qualifiedly guess could still be estimated.

In chapter 3 six methods of gap filling was evaluated. It turned out that the simple linear method performed quite well in comparison to the others. The envelop method did also perform acceptable for smaller gaps. Therefore are these two method evaluated on the ECO dataset, to see if it is possible to increase performance when preforming gap fixing.




Figure 4.10: Average recognition after recovery of the 3 methods

Like the error tolerance experiment described in section 4.5.2, errors are artificially introduced in the dataset. After this process, an attempt is made to remove the errors by using gap filling techniques.

On figure 4.10, the average F1 score is shown for all 3 recognition methods applied on all 5 appliances. This indicates that by using a simple reconstruction method, it is possible to keep the high performance. Where for the more advanced Envelope method, it seems to make it worse. This is typically because the more advanced types of reconstruction methods try to model the frequency components, which for greatly stochastic signals is more likely to introduce errors.

To ensure that it is not one recognition method that causes the bad performance for the envelope, the results are inspected for the individual appliances and methods.




Figure 4.11: Recovery in FHMM algorithm

On figure 4.11 is the result from the reconstruction with the FHMM algorithm shown. Here we see that the performance is prevented from decaying by the two reconstruction methods. It is also shown that the Envelope reconstruction actually preforms marginally better for the FHMM algorithm. In all cases is the reconstruction better than non-reconstruction.




Figure 4.12: Recovery in Parson algorithm

On figure 4.12 is the result from the reconstruction with the Parson algorithm shown. Here we see that the Envelope method is always worse than no reconstruction. For some of the appliances is the linear reconstruction a help, but fore some does it preform worse than no reconstruction.

The parson algorithm looks at jumps in the power, and uses a moving window mean filter to noise reduction and to ensure that only the significant jumps is detected. This behaviour can explains why a linear reconstruction gives almost no performance gain compared to no reconstruction for some of the devices.




Figure 4.13: Recovery in Weiss algorithm

On figure 4.13 is the result from the reconstruction with the Weiss algorithm shown. For the Weiss algorithm it is appliance depended what reconstruction method is the best. Even though it seems like reconstruction is not preferred for some of the appliances like the laptop, is it worth noticing that the score is close to 0.2 which is a low score. For appliances that have a higher score like the TV and fridge reconstruction seems to work as intended.

The low score indicates that the method in it self is struggling to identify the appliance, and the data error is therefore not the primary problem.

4.6 Chapter Discussion

In the chapter the effects of errors and gaps in the dataset is investigated. It is shown that a system that simply ignores that errors will perform worse than one that tries to do some compensation. The amount and type of compensation is shown to be very signal depended. In general is a simple linear interpolation enough when using algorithms that are designed for low frequency signals.

The resolution of the samples seems to have an effect of the performance. The algorithms in focus are designed to low frequency signals, even so it is shown that the small variations created in higher resolution can improve the performance. This is working with the assumption that the instantaneous power consumption received from the smart meter is a average of the power consumption since the last sample.

There is also indications of that some appliances is more easily detected by some algorithms, and other appliances by other algorithms. This indicates that it is very likely that a greater performance can be achieved by mixing algorithms depended on the appliances.

Environment Influence 5

Not just the sample rate and error rate are determining factors when designing a non-intrusive load monitoring (NILM) application. The "environment" which the application is deployed and trained in is critical for the performance of the system. The "environment" is the parameters describing the static conditions of the household which the algorithms are deployed in. An example of an environment parameter can be number of known devices, number of unknown devices, number of simultaneously active devices or number of training days.

In order to investigate some of the environment parameters effect on a NILM application the SmartHG dataset is used. In the chapter is appliance disaggregation conducted on the SmartHG dataset, using the factorial hidden Markov models (FHMM) and Parson algorithms. The Weiss algorithm uses the power factor information, which is not available in the SmartHG dataset. The disaggregation process consists of a training phase and a validation phase.

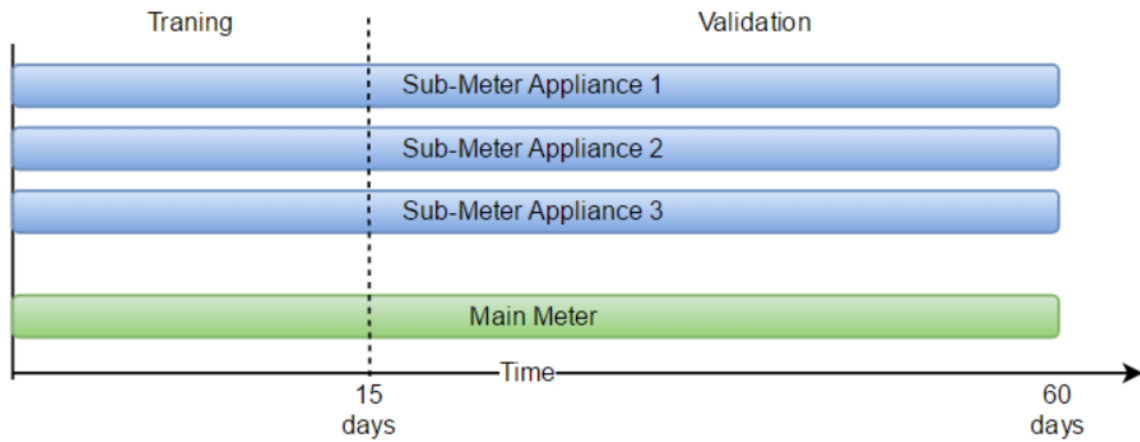


Figure 5.1: Disaggregation phases

All experiments conducted in the chapter will start with a training phase of 15 days. During this period the sub-meter and main meter data will be used to by the FHMM and Parson algorithm to learn statistical disaggregation models. After this initial training phase will the algorithms be tasked with the job of disaggregating the main meter signal in a 45 day validation period. The algorithms will only have access to the data supplied by the main meter, and the statistical disaggregation models found in the initial training phase. The sub-meter data in the validation phase be used to validate the results of the disaggregation. The basis setup is illustrated in figure 5.1 where the green color symbolises data from main meter, and blue symbolises data

from sub-meter.

5.1 Challenges In The SmartHG Dataset

The Electricity Consumption and Occupancy (ECO) dataset used in the last chapter, consists of 6 households and the data is sampled at a rate of 1 Hz over a period of 8 months. The SmartHG dataset is different in many aspects. The data is sampled at a slower rate of $\frac{1}{30}$ Hz, but on 25 households over a period of 8 month. Each house have only a small number of sub-meters. The sub-meters are mainly placed on little consumers such as televisions and stereos, which presents an interesting challenge for load disaggregation. The SmartHG dataset contains both the accumulated power, and the instantaneous power usages of the different households.




Figure 5.2: Frequency comparison of the reconstruction methods

The power distribution of three households is shown in figure 5.2. This illustrates how much energy each appliance uses, in relation to the households total energy consumption. The "other" category shows the energy consumption not accounted for by the sub-meters.

Household 3,10 and 18 are some of the houses that have the smallest *other* category. They are therefore selected for further study, since they provide the most information about the house.

5.2 Background Consumption Influence On Disaggregation

The households in the SmartHG dataset have a relative big consumption created by appliances in the *other* category. The consumption in the *other* category is created by appliances that are unknown. In order to do disaggregation, statistical models like models based on the FHMM and Parson algorithms, must first be created. This requires knowledge about the appliances consumption. This kind of information is not available for appliances in the *other* category.

This consumption created by appliances that is not part of the statistical disaggregation models can be thought of as "background consumption". The appliances in the *other* category will always create *background consumption*. This does not exclude appliances that are not part of the *other* category from creating *background consumption*. All appliances that are not part of the statistical disaggregation models, contribute to the *background consumption*.

5.2.1 Detection In An Environment With High Background Consumption

To investigate the influence *background consumption* have on a NILM application, disaggregation of the known appliances of house 3, 10 and 18 are done using the Parson and FHMM algorithms.

	FHMM		Parson	
	F1	Accuracy	F1	Accuracy
TV 1	0.19	0.74	0.10	0.40
PC	0.19	0.84	0.13	0.45
TV 2	0.03	0.84	0.20	0.11
Average House 3	0.14	0.81	0.14	0.32
TV 1	0.60	0.76	0.40	0.25
Stereo	-	1.00	-	1.00
PC	-	0.99	-	0.99
TV 2	-	0.99	-	0.99
Average House 10	0.15	0.94	0.10	0.81
TV 1	0.36	0.65	0.24	0.14
Lamp	-	0.99	-	-
TV 2	0.73	0.58	-	0.00
Average House 18	0.36	0.74	0.08	0.04

Table 5.1: Appliance disaggregation results of house 3,10 and 18 for the smartHG dataset

As previously stated is the *background consumption* high for this dataset. Table 5.1 summarizes the results from the disaggregation. The low F1 scores indicates that it is hard for the algorithms to correctly disaggregate the meters. This is due to the many spikes there are in the data that can indicate a change in appliance usage. In figure 5.3 is a snippet from the validation period shown. The blue lines shows the data as it is received from the smart-meter. The other lines indicates the appliances. In the sub plot named "True Signals", on the left side, is the true consumption shown for all the meters. This consumption is known from the sub-meters on the appliances. In the sub plot named "Inferred Signal", on the right, side is the inferred usages of the different appliances created by the FHMM algorithm.




Figure 5.3: FHMM disaggregation snippet

In the figure does the blue lines illustrates how the main meter have many fluctuations. The

algorithm tries to map each fluctuation to a change in an appliance state, or as *background consumption*. There are many of the fluctuations created by the *background consumption* that are similar to the power changes created by the TV, PC and entertainment system. The inferred signal shows how the algorithm thinks the appliances is used. It is easily seen that this is quite different from the true values.

5.2.2 Detection In An Environment With No Background Consumption

The assumption that the *background consumption* is what troubles the disaggregation dictates that for a house where all appliances are known, and there are no *background consumption*, the disaggregation will perform better.

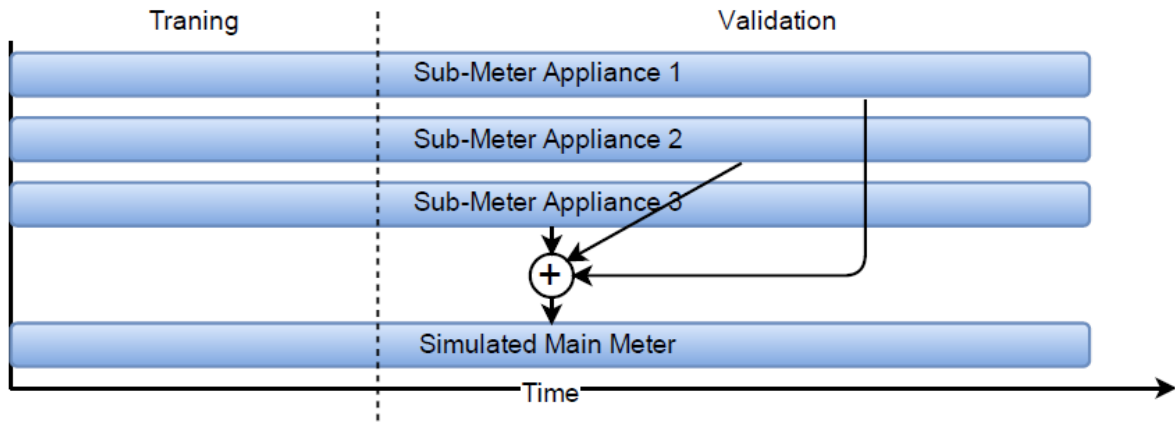


Figure 5.4: Artificially constructed main meter

In order to test this hypothesis is a main meter signal artificially constructed by summing the values of all the sub-meters in a house. This process is illustrated in figure 5.4. In this manner is a new set of artificial houses created where there are no *background consumption*. By applying the same disaggregation algorithms, as for the environment with high *background consumption*, on the new environment with no *background consumption* is the performance much increased as shown in table 5.2.

	FHMM		Parson	
	F1	Accuracy	F1	Accuracy
TV 1	0.73	0.96	0.26	0.86
PC	0.74	0.96	0.30	0.91
TV 2	0.83	0.96	0.20	0.11
Average House 3	0.77	0.96	0.26	0.62
TV 1	0.97	0.98	0.40	0.25
Stereo	-	1.00	-	1.00
PC	-	0.99	-	0.99
TV 2	-	0.99	-	0.99
Average House 10	0.24	0.99	0.10	0.81
TV 1	0.95	0.98	0.24	0.14
Lamp	-	0.99	-	-
TV 2	0.73	0.58	0.73	0.58
Average House 18	0.56	0.85	0.32	0.24

Table 5.2: Disaggregation of appliances on artificially constructed main meters

The F1 and accuracy score are both improve, as shown in table 5.2. Even though the artificial house only contains three or four appliances, and all are accounted for by the statistical disaggregation model created by the FHMM and Parson algorithms, there are still some of the appliances that are hard to find. This is due to the appliances power usage pattern in OFF/standby mode combined with the rare usage of the devices. When the standby pattern of a device is complex and the appliance is only rarely used the model will try to fit to the standby pattern and not the usage pattern. This happens if the over all best consumption disaggregation results is achieved by fitting to the small energy fluctuations in the OFF state. This is often the case for appliances that are rarely used int the training period. The downside to this is that it will be incapable of detecting change between the ON/OFF states.

One other aspect to note is that even though the F1 scores are higher, only a few is higher than 0.9. This could be due to "appliance interference". This happens when appliances are similar. When a house contains two or more appliances that are similar it gets harder for the algorithm to tell if it is the one or the other. This is could be what happened for TV 1 and TV 2 in house 3. The algorithms will wrongfully assign some of the events belonging to TV 1, to TV 2 and vice versa. This happens because the power draw signature are similar. This can be a problem if the goal is to determine exactly which appliance there is responsible for the power consumption. This is one of the key aspects of the little consumers problem, discussed in Chapter 4.

Appliance interference has a profound effect on the Parson algorithm. The parson algorithm is designed to split up appliances based on appliance categories, by utilizing a set of general models. Since a lot of the appliances are similar in their consumption changes from state to state are the model used by the Parson algorithm almost identical which lead to the events being wrongly categorised. The usages general models is one of the benefits of the Parson algorithm, since it makes the algorithm very portable between unknown domains. Unfortunately does this feature come whit the downside of increased sensibility to *appliance interference*. The influence of *appliance interference* is further studied in Section 5.3.2.

5.2.3 Background Consumption Effect On Training And Validation

The experiments in the two previous Sections 5.2.1 and 5.2.2 suggest that *background consumption* has an impact on the performance of the system. But for a system based on machine learning techniques like the FHMM and the Parson algorithms it raises the questions, is it too hard for the system to train the correct statistical disaggregation models in the signal with much *background consumption*, or are the *background consumption* just interfering with the disaggregation process?

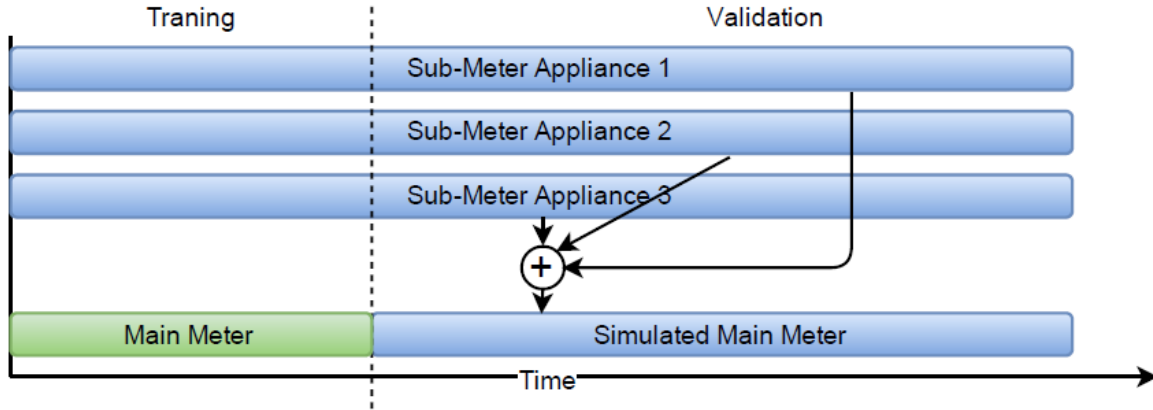


Figure 5.5: Illustration of dataset creation by combining real and constructed data

In order to investigate these questions a new dataset was created that combined real data collected from the house main meter, and an artificially constructed main meter. The new dataset was created by using the real data in the training period of the house, and constructed main meter data in the validation, as illustrated on figure 5.5.

This creates an environment with a lot of *background consumption* in the training phase and no *background consumption* in the validation phase. If the *background consumption* did not influence the creation of the statistical disaggregation models that represent each appliance, then the results were expected to be just as good as for the noise free experiments conducted in Section 5.2.2.

	FHMM		Parson	
	F1	Accuracy	F1	Accuracy
TV 1	0.02	0.94	0.23	0.86
PC	0.29	0.94	0.32	0.92
TV 2	0.00	0.88	0.20	0.11
Average House 3	0.11	0.92	0.25	0.63
TV 1	0.00	0.74	0.40	0.25
Stereo	-	1.00	-	1.00
PC	-	0.99	-	0.99
TV 2	-	0.99	-	0.99
Average House 10	0.00	0.93	0.10	0.81
TV 1	0.00	0.86	0.24	0.14
Lamp	-	0.99	-	-
TV 2	0.73	0.58	0.73	0.58
Average House 18	0.24	0.81	0.32	0.24

Table 5.3: Disaggregation in real and constructed main meters combined

If the results from the experiment, shown in table 5.3, is compared with the results from table 5.2 from section 5.2.2 it is clearly seen that the performance of the system trained in real data is much lower than the one trained on artificially constructed data. This indicates that the disaggregation models obtained in the real data is influenced by the *background consumption*, and is therefore not fitting accurately enough to the true disaggregation models. This indicates that the low performance in the real data shown in table 5.1 in section 5.2.1 is caused by the models not fitting correctly to the true values, and not by the models being triggered by application noise from other appliances.

In order to validate this hypothesis was a similar experiment conducted, where the training data was artificially constructed, and the validation data was the real data from the main meter as illustrated in figure 5.6.

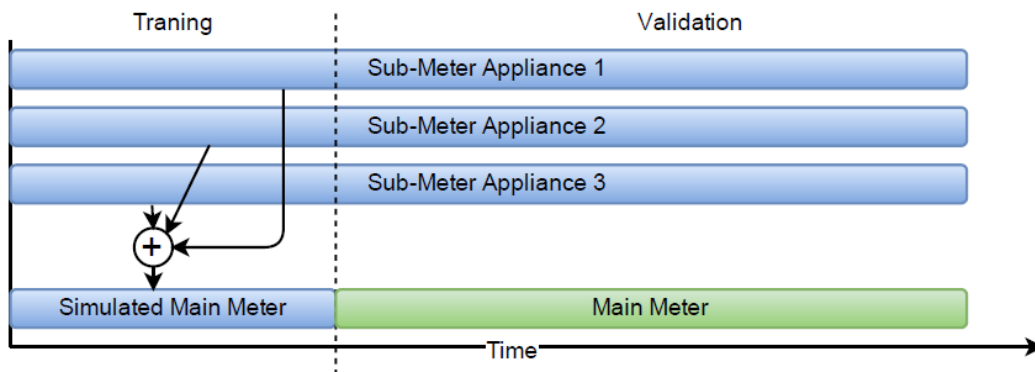


Figure 5.6: Illustration of dataset creation by combining constructed and real data

By training the model on the constructed data, is the true model for the appliances more easily found. If the environment with high *background consumption* and real data is non-interfering with

the disaggregation models can a performance like for the artificially constructed environment, with no *background consumption*, described in Section 5.2.2 be expected for the validation phase.

	FHMM		Parson	
	F1	Accuracy	F1	Accuracy
TV 1	0.10	0.42	0.10	0.45
PC	0.13	0.39	0.12	0.44
TV 2	0.21	0.35	0.20	0.11
Average House 3	0.15	0.39	0.14	0.33
TV 1	0.40	0.25	0.40	0.25
Stereo	-	1.00	-	1.00
PC	-	0.99	-	0.99
TV 2	-	0.99	-	0.99
Average House 10	0.10	0.81	0.10	0.81
TV 1	0.27	0.26	0.24	0.14
Lamp	-	0.99	-	-
TV 2	0.73	0.58	0.73	0.58
Average House 18	0.33	0.61	0.32	0.24

Table 5.4: Disaggregation in constructed and real main meters combined

The results are shown in table 5.4. The performance of this experiment is lower than for the one in the environment with no *background consumption*. This lead to the conclusion that the *background consumption* is both affecting the models, and interfering in the validation process.

The results from the disaggregation of the data in the real environment from section 5.2.1 are compared with the results from the the previous experiments, where the models were extracted from constructed data and validated on the real data. The results is summarized in table 5.5.

	FHMM Real		FHMM Constructed	
	F1	Accuracy	F1	Accuracy
TV 1	0.19	0.74	0.10	0.42
PC	0.19	0.84	0.13	0.39
TV 2	0.03	0.84	0.21	0.35
Average House 3	0.14	0.81	0.15	0.39
TV 1	0.60	0.76	0.40	0.25
Stereo	-	1.00	-	1.00
PC	-	0.99	-	0.99
TV 2	-	0.99	-	0.99
Average House 10	0.15	0.94	0.10	0.81
TV 1	0.36	0.65	0.27	0.26
Lamp	-	0.99	-	0.99
TV 2	0.73	0.58	0.73	0.58
Average House 18	0.36	0.74	0.33	0.61

Table 5.5: Trained in real vs. constructed data

The table shows that the performance of the system is better when trained in the real environment, as in relation to when the the system is trained in a artificial environment and then deployed in a environment with *background consumption*. This is in contrast to what one might think, since the environment with no *background consumption* should have supplied more correct models. But when the *background consumption* is removed from a house, in the training process, is a offset error introduced in the models.

Appliances such as refrigerators and freezers that are always on will create a local house offset. This will vary from house to house, and is depended on the types and number of appliances in the house. If only a small number of devices are modelled in each house, as in the SmartHG case, is the offset almost exclusivity a part of the *background consumption*. When the models are learned in the training phase in the real data, is the offset learned into the model, and further improves the model. This is not the case when the model are trained in the constructed dataset.

If not all appliances in a house are known, a better performance can therefore be obtained by training in the intended environment of deployment.

Some algorithms are designed to not be affected by the house offset, and is therefore moved from different environments more easily. The Parson algorithm is designed with this in mind. The performance of the Parson algorithm is therefore the same, or in some cases improved, when using models trained in constructed data. This advantage unfortunately comes with the problem of generalization, which lead to harder source separation, as discussed in section 5.2.1.

5.2.4 Improving Performance Using Norm Knowledge

When an appliance is in a environment with a lot of *background consumption* some of the *background consumption* can be interpreted as the appliance. This *appliance interference* creates a scenario where it seems like the appliance is used more than it actually is. Often does the *appliance interference* create usage patterns that is in contrast to that one would expect. This can make it seems like a appliance is used in a unrealistic manner. By improving the method with the knowledge of how a specific appliance is most commonly used, also called the norm usage, it is possible to filter the signal and improve the performance.

In figure 5.7 is the disaggregation of TV 1 from house 10 shown for an one day period. The blue line indicate the FHMM algorithms suggestions on when the TV are ON and OFF. The green line is the actual ON and OFF periods, collected from the sub-meters.




Figure 5.7: Disaggregation of TV 1 from house 10

As seen in the figure there is a lot of small spikes that are less than 10 min, and sometime is the TV ON for 15 minutes off in 2 and then on again for another 15 minutes. This behaviour does not fit the normal behaviour one would expect for a person. The norm is to see TV more than 10 minutes of the time, and when a TV is turned off it is common that it is turned off for atleast 15 minutes before turned on again. These numbers are a loose estimate, found by observing the sub-meter data.

By applying purge and merge process, where all events shorter than 10 minutes are purged from the signal and events separated with less than 15 minutes of off time are merged, it is possible to clean the signal.




Figure 5.8: Purged and merged signal from TV 1

In figure 5.8 is the same signal as in figure 5.7 filtered by using the purge and merge process. The purge and merge process can be used to create a norm filter, that filters appliance usage that does not fit the norm. As seen in figure 5.8 does the signal inferred by the filtered FHMM algorithm match more closely to the actual signal.

In this example is the norm knowledge applied as a filter, that can filter the disaggregation result. Another approach could be to integrate the norm knowledge as a component in the state transition probability of the hidden Markov models (HMM).

5.3 Model Size And Completeness Influence

Some of the parameters that often differ in the many different environments NILM applications are deployed in, is the number of appliances in the environment, and the number of appliances that are known in the environment.

The environment size or complexity is a parameter that tells how many devices that are in a given environment. As the complexity in a NILM application is increasing the detection rate is decreasing [28]. This is one of the major problems and is why many researchers only focus on a small subset of appliances that have a fairly unique consumption signature.

The model completeness is a metric describing how many appliances are known by the application in relation to the total amount of appliances. This can also be seen as the amount of appliance noise as discussed in section 5.2.

5.3.1 Test Set Creation

In order to experiment with the completeness and complexity in the SmartHG dataset is a more controllable series of datasets artificially constructed from the SmartHG dataset.

First an artificial house where created called the "TV House" dataset. The TV house dataset is created by picking the relative dominant TV signals from the houses 5,10,11,13,18 and 23 in the SmartHG dataset and combining them to one artificial house with 7 TV's.

This dataset contains of dominant appliances, and there is no appliance noise from unknown devices. It is still worth noticing that all the 7 appliances are a Type-II appliances, and they are all TV's, which make there usage pattern some what similar. Since the Parson algorithm is greatly effected by this will the focus mainly be on the FHMM algorithm.

If the TV House dataset was seen as a set of sub-meters as mathematically shown in equation 5.1.

$$H_{full} = \{TV_1, TV_2, ..., TV_7\} \quad (5.1)$$


Since there is no appliance noise, can the main meter data be found as the sum of sub-meters data as shown in equation 5.2.

$$M_{main} = \sum_{X \in H_{full}} X \quad (5.2)$$

Using this information it is possible to create a set with a specified complexity, by taking a subset of the H_{full} set with the cardinality of the specified complexity. The main meter can now be found on this subset using the same principle as in equation 5.2. Using the equation 5.3, it is possible to find a set of all sets with a desired complexity c .

$$\mathbf{H}_c = \{x | x \in \mathbb{P}(H_{full}), \forall |x| = c\} \quad (5.3)$$

This can be graphically shown as in figure 5.9.

Figure 5.9: \mathbf{H}_c sets from the TV house dataset

Where the \mathbf{H}_7 only contains the H_{full} set. In this manner it is possible to artificially create one or more datasets in a desired complexity equal to or less than the H_{full} complexity. The amount of artificiality created houses for a given complexity can be found by using the simple combinatorial equation shown in equation 5.4.

$$|\mathbf{H}_c| = \frac{|H_{full}|!}{c! \times (|H_{full}| - c)!} \quad (5.4)$$

$$A_c = |\mathbf{H}_c| \times \frac{c}{|H_{full}|} \quad (5.5)$$

As illustrated in figure 5.9 can an appliance appear in multiple sets in a given \mathbf{H}_c collection. The number of sets containing a specific appliance in a \mathbf{H}_c collection A_c can be calculated as in equation 5.5. In order to ensure that the combination of appliances does not have an effect on the experiments, are the experiments conducted on all combinations in \mathbf{H}_c . The results for each experiment is an average of the performance for the appliance in the experiment.

	$c = 1$	$c = 2$	$c = 3$	$c = 4$	$c = 5$	$c = 6$	$c = 7$
$ \mathbf{H}_c $	7	21	35	35	21	7	1
A_c	1	6	15	20	15	6	1

Table 5.6: TV House dataset complexity

In the case of the TV house dataset can the $|\mathbf{H}_c|$ and A_c values for the different complexity be seen in table 5.6.

5.3.2 Model Complexity Test

To investigate the effect the number of appliances in a house have on the performance of a NILM system, is disaggregation done on every artificial house in \mathbf{H} .




Figure 5.10: Average score at different complexity

The test is performed on every combination, to ensure that there is not some combination that are more favourable than others. The results shown in figure 5.10 are the average of the results from the many combinations as described in section 5.3.1.

For the most part it is fairly easy to correctly classify the TV's when the complexity is only one. As a general trend does the F1 and accuracy score of the TV's decrease as the complexity of the house increase. This is due to the application spill over effect. The spill over effect is when there actually is an event, but the system classifies it to the wrong device. An effect of this is most clearly seen on TV 2 signal. TV 2 is a very complex signal, and is therefore hard to track by the algorithm, as seen at the \mathbf{H}_1 complexity that have a F1 score on almost zero. When the complexity is increased it looks like it is easier to detect the signal, and the F1 score is increasing. What is actually happening is the spill effect from the other TV's. Since all the appliances are TV's they have a lot of overlap in their usage. A lot of people watches TV at the same time of day e.g. between 18-22. This makes the algorithm encounter signal's that are similar in structure for the different TV's, and it have a hard time deciding which appliance is responsible. Therefore can some of the events generated by the other TV's spill over in TV 2. If by chance TV 2 actually was on when the events from the other TV's was wrongfully assigned, the F1 and accuracy score of TV 2 would improve.

For other appliances that are not as hard to classify as TV 2 will the spill over effect decrease the F1 and accuracy score as seen on the figure.

5.3.3 Model Completeness Influence

In the experiments conducted on the the SmartHG dataset in section 5.2 it was shown that the appliance noise made it hard to correctly disaggregate the appliances of the houses. In the last experiment in section 5.3.2, it was shown that the performance of the disaggregation was higher when the complexity of the house was low. This is much the same results as seen in the experiments on the SmartHG dataset, where the completely simulated house, that had a low complexity, preformed better than the real house with a high complexity.

But in the SmartHG dataset was the complexity of the real dataset unknown since the amount of appliances that generated the appliance noise is unknown. It is therefore hard to conclude if the poor performance in the real data is due to the fact that the appliances is unknown by the disaggregation model, or because the complexity of the house is greater for the real data.

To further investigate this was an experiment conducted on the TV house dataset, where the house models where trained on the \mathbf{H}_7 dataset. The models trained did not contain the full knowledge of the house. The models was only trained to find a subset of meters in the house. Like for the experiments in section 5.3.2, was the experiment conducted on all subset combinations.




Figure 5.11: Completeness test of disaggregation of appliances in the TV house dataset

In figure 5.11 is the result for each appliance shown. The result is an average of all tests at a given model completeness. It is seen from the figure that even though the disaggregation algorithm obtains more information about the complete house, it does not improve the recondition.

This indicates that the 7 algorithms working in parallel, each disaggregating one meter, would obtain the same results as one model capable of seeing the house as a hole. It is worth mentioning that these tests are based on the FHMM algorithm and other algorithms might preform different in this aspect.

The results indicate that the complexity of the house is a more determining factor for the performance than the completeness of the model.

5.4 Chapter Discussion

The findings in this chapter indicates that the complexity of a dataset has a profound impact on the detection performance. A simple way to decrease the performance in a house is to look at each phase in the house separately. This is one of the techniques used in the ECO dataset to improve the performance in the system. This does require additional knowledge of which appliances that are connected to which phases. This knowledge can be analytical extracted from the dataset, or must be supplied at training time.

It is also indicated that the methods based on the HMM does not greatly improve by knowing the entire house model, in relation to only knowing a sub part. This can indicate that a system that focus on detecting each appliance individuality, potentially can detect the appliance as well as one that tries to see the house as a hole.

The many problems that can occur when using small consumers and once that are similar in behaviour is nicely illustrated in the chapter. The spill over effect for similar appliances is greatly decreasing the overall performance, and the small consumers is almost undetectable. In order to better this performance, more unique features is needed. This could either be higher frequency features, as discussed in chapter 4, or by using reactive features. Furthermore it is shown that the best performance is obtained by training in the same environment as the solution will be deployed in. This is a subject that is also discussed by the developers of the Parson and Weiss algorithms [3] [25], where they discussed ways to train or fit the models to the desired deployment environment.

NILM As An Application 6

The last couple of years have the number of smart meters installed in residential houses increased drastically [ref from book](#). The motivation for installing smart meters in the different houses have been to better understand energy consumption, in order to better plan energy distribution and production [ref from book](#). With the smart meter came a golden opportunity for non-intrusive load monitoring (NILM), since the equipment and infrastructure needed to measure the consumption and transfer the results through internet was available in many households.

The applications proposed in many of the articles published about NILM focuses on energy management. Either it is for the electricity producers, that needs it to better predict consumption, or it is for the resident of the house that can optimize power usage to obtain savings [1]. Even though these are fine examples of the potential usages of NILM there might also exist an opportunity for the electric companies to sell the NILM information and act as a data broker. This chapter contains a small case study to illustrate this point.

6.1 The TV Viewing Habits Case

The case setting is a small town, that have a number of households and an electric company supplying energy to the city. For this case data from the SmartHG dataset is used to simulate a small town with 4 houses. The houses selected for this case are 10, 13, 18 and 23 since they contain a TV that are relatively dominant, as discussed in section 2.2.2.




Figure 6.1: The city setup

The city setup is illustrated in figure 6.1. The figure illustrates how the electric company supplies energy to the houses. All the houses have smart meters, and are informing the power company about the current consumption of each house, using some network. This is a fairly simplified

example of a smart grid, that exists in many cities today. In this case it is assumed that the electric company receives information about the consumption of the houses each 30 seconds, as in the SmartHG dataset. Even though many smart meters today are capable of delivering information at this speed, it is more common to only get information each hour. This is mainly because the information is used to regulate distribution and production, which is a slow process that would not improve by faster update rates.

The electric companies collect the information, and by using NILM can calculate statistical information about the town. In the fictive case is a local TV station interested in knowing what time a day the citizens of the town watches TV, in a specific period, in order to improve some aspect of there product.

6.2 Methodology

The aim of this case is to illustrate the potential of this kind of application for an electric company. Therefore is the data analysis conducted in a manner, that illustrates how an electric company could do it. This means that only the main meter information is used in the analysis, and the sub-meter information is only used to validate the results. The goal is for the electric company to disaggregate the TV signals from the main meter signal, in order to use it in some statistical analysis.

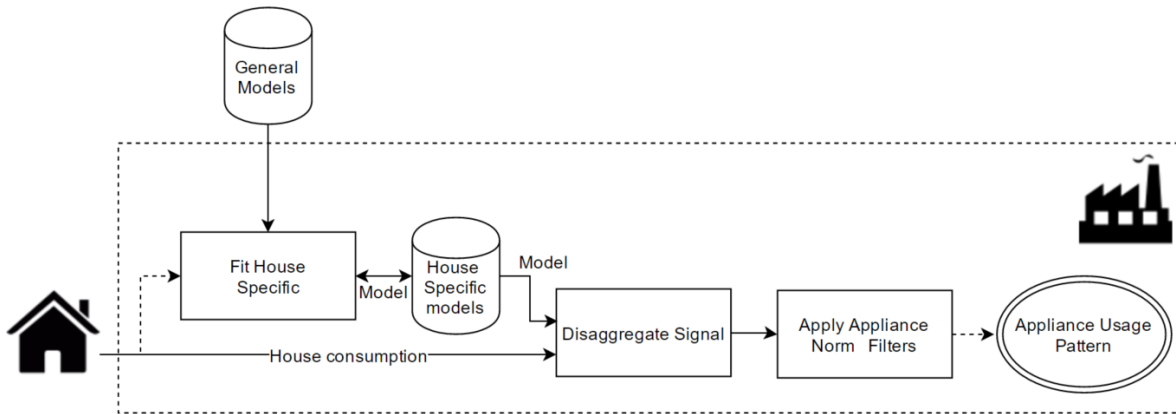


Figure 6.2: The disaggregation process

This process is illustrated in figure 6.2. It is assumed that the power company have access to a very general statistical model of the TV. This model can come from some outside database, or be one that the electric company developed themselves. The first step is to create a household specific model, that is a specialisation of the general model. This could be done with sub-meter data, but to make the experiment more realistic is the fitting of the model done using only the aggregated data from the main meter. This is done using the same method as developed for the Parson algorithm [3]. This method finds periods where an appliance is turned ON and OFF without any other appliances changing states, is detected. This on/off single event fragments are then mapped to specific appliances using the general model. When enough on/off single event fragments have been collected can this information be used to fit a more specific model. This

can happen in a contentiously manner, where the specific model is improved as more data is collected.

Using the specific models disaggregation of the main meter signal can now be achieved. The statistical models selected in this case are the factorial hidden Markov models (FHMM). As discussed in section 5.2.1, the models alone gives a very error prone signal in the FHMM. The signal is therefore feed through a norm filter. As discussed in 5.2.4 can the information about how a user normally would use the appliance be used to filter the disaggregated data to obtain better results.

The result of this process is an appliance usage pattern, that can be used by the electric company to derive various statistical properties.

The experiment is conducted with four weeks of data been used to fit TV specific models for each house. The analysis for the TV station is conducted over a period of six weeks, that is not overlapping with the initial four training weeks.

6.3 Results

The viewing habits of the small town obtained by the disaggregation analysis is compared with the actual viewing habit obtained from the sub-meter data. In figure 6.3 is the results shown for the first week.




Figure 6.3: Viewership in week one

The figure shows the viewership, that tells how many in the city is watching TV at a given time. The green line is the actual viewership, where the blue is the viewership reported by the electric company. The results differer, as one would expect, since the disaggregation process in a noisy environment is a hard task. The general trends that many are watching TV in the evening, and not that many in the morning, is still represented nicely. This illustrates that the general trends of the viewership is maintained in the disaggregated signal, and this is commonly what is important for this kind of costumer statistics.

In table 6.1 is the average number of viewers in 3 hour timeslots shown. The results shown in the table is the average of all 6 weeks, split up in the different weekdays.

		Time of day							
		00-03	03-06	06-09	09-12	12-13	15-18	18-21	21-00
Monday	Inferred	0.12	0.00	0.66	1.19	1.02	1.31	2.39	2.08
	True	0.13	0.28	0.98	0.87	0.68	1.28	2.49	1.72
Tuesday	Inferred	0.13	0.28	0.98	0.87	0.68	1.28	2.49	1.72
	True	0.19	0.09	0.39	0.88	0.73	1.24	2.23	2.26
Wednesday	Inferred	0.19	0.09	0.39	0.88	0.73	1.24	2.23	2.26
	True	0.09	0.30	0.58	0.79	0.77	1.42	2.72	2.11
Thursday	Inferred	0.09	0.30	0.58	0.79	0.77	1.42	2.72	2.11
	True	0.12	0.09	0.31	0.96	1.12	1.69	2.86	2.30
Friday	Inferred	0.12	0.09	0.31	0.96	1.12	1.69	2.86	2.30
	True	0.10	0.18	0.66	0.91	0.74	1.65	3.17	2.35
Saturday	Inferred	0.10	0.18	0.66	0.91	0.74	1.65	3.17	2.35
	True	0.28	0.17	0.95	0.77	1.39	1.36	2.26	2.10
Sunday	Inferred	0.28	0.17	0.95	0.77	1.39	1.36	2.26	2.10
	True	0.12	0.00	1.01	0.96	0.96	1.33	1.93	1.65

Table 6.1: Average viewership in a 6 week period

This illustrates that the true trends in the information is maintained. The results indicates that it is in some degree possible to obtain user information from the disaggregated data. The TV is a relative hard appliance to detect, since there are many different types. The success in this experiment is partly due to the fact that the TV was some of the main consumers of the houses selected, and they had a pure Type-I behaviour.

6.4 Chapter Discussion

In this chapter was a very simple example of how the electric company can extend their product range to statistical information. This information could potentially also be used to extend their product range to their current costumers. One example of this could be a fridge surveillance system that would send warning messages to the resident if their fridge stopped working.

At the current state of NILM is the information only limited to a small amount of devices, since it is hard other appliances than the top consumers in the household. One of the techniques shown to greatly improve the effectiveness of NILM is to switch from the sub 1 Hz sampling range to a higher sampling rate in the kHz range. In the future where either meters capable of delivering information at high frequency is common in the houses or the NILM techniques have advanced, it is a possibility for the electric companies to deliver very detailed user behavioural statistics. This information can be of interest for companies developing and selling electronics, since they are able to get accurate user reports about their products.

Discussion 7

Non-intrusive load monitoring (NILM) is an area of research that is gaining a lot of attention due to the more complex smart grid structures. Some of the challenges in the smart grid is that there is not just one producer and a series of consumers. There can be multiple of both, and some households contains solar panels, or other energy producing technologies, that makes them both producer and consumer at different times. This creates challenges for some of the NILM applications, as it can camouflage the true power consumption of the houses. The same problem can be created by battery devices, that can mask the a appliance power demand. Some of these problems can be observed by monitoring the data quality.

7.1 Creating High Quality Measurements

When monitoring the the data quality as in chapter 2 a range of problems when collecting consumption data can be detected. When observing the collection of SmartHG data there were two error gropes that reviled itself. The one was equipment failure. Both failure for a specific meter, or the server was detected. The second is the houses that exhibit masking behaviour. By looking at the information collected by the sub-meters and comparing it with the main meter data it is possible to detect the masking behaviour created by devices such as solar cells or backup batteries.

By detecting the errors during the measurement period it is possible to correct the errors. This creates better and more clean measurements. By looking at the measurements activity, it is possible to find still and active areas. This can be used to find active and inactive periods of the devices. These periods can be used to create a more direct training of the statistical models used to do appliance disaggregation. Finding periods in the data where only one single appliance have a ON/OFF cycle, without any other appliance changes state, have shown to be an important training technique. This technique is mostly relevant if sub-meter information is not available. Utilizing the information obtained about the activity in the quality analysis can help find such spots.

One of the big challenges in a NILM application is the amount of devices that is operating on the same power line. A model that needs to disaggregate a signal containing many devices preforms worse than one where there only is few devices. It is therefore beneficial to look at the power usage of each phases individually. By using this separation can the complexity of the models be reduced greatly, which is shown to increase performance.

Having these points in mind when measuring data, at the consumer, can potentially create a better dataset. A dataset could also improve by supplying information about what devices is not measured in a research study. If a list of devices in the household was supplied with

the measurements it is possible to better gesture about the background consumption created by the other devices. This kind of information would also be beneficial for methods based on unsupervised learning [23].

7.2 Capability Of NILM

In cases focused around selected appliances have NILM shown to be quite effective. But when seeing NILM as a tool to disaggregate a whole household appliances, there is still many challenges. Generally it is easy to detect the top consumers in a household. The big consumption increases the chances of a unique power draw pattern. The more unique power draw pattern the easier it is to detect the appliance. The majority of devices in a household only consume a small amount of the total energy. These devices is hard to disaggregate. This is because the power draw pattern for the devices is relatively similar. This causes the device disaggregation models to interfere with each other under the disaggregation process. The result indicates that sometimes some of the events from one device is classified to the wrong devices.

Much research have been focused on solving this problem. The general concept behind all the solution is to find other features than the static power draw that can create uniqueness. One approach is to use the the difference between the voltage and current usage of the devices, known as the power factor. The power factor can indicate if a load is inductive, capacitive or resistive in nature. One other approach is to sample at high rates in the kHz range. This can supply information about transients, harmonics or noise that can be used to detect uniqueness.

The solutions most commonly used today focus on techniques that uses low sample rates, in the sub 1 Hz range. This is due to the fact that most smart-meters are capable of delivering data at these rates. A new type of meter must be developed in order to utilize the benefits that higher sample rates would create. This meter must either sample or transmit the data at higher rates, or process the data on sight and transmit the results. This is a costly process, since a new meter must be developed and installed in the households. By including this type of capabilities in the next generation of smart-meters it would be possible to advance the NILM research area. This would increase the businesses opportunities for consumer applications that utilizes NILM techniques.

7.3 NILM Deployment

In the current state can NILM only deliver a very general picture of the top consumers. What type of devices that is the top consumers is different from country to country. In countries that get hot in the summer, is the air condition almost at the top of the list. In countries that are cold and have a rocky underground uses a lot of power on electric heading. The difference in consumption culture around the world makes it hard to cross compare results between datasets. This makes it hard to determine which solutions is most suited for a desired country.

When validating only by looking at the F1 and accuracy scores the results seems promising. When a appliance have a 0.60 F1 score, it generally means that the ON periods was guessed correctly 60% of the time. This still leaves 40% of the ON time to be wrongly allocated. The wrong allocations can be both be ON events that are missed and ON events wrongly introduced.

This gives the same F1 score, but can make a difference in a application. This is illustrated in 7.1, where two versions of the same event signal is shown.




Figure 7.1: TV event detection

The "Raw Disaggregation" is before applying a simple norm filter, the second illustration is after the norm filter. The norm filter used is the same as in in Section 5.2.4. Visually the difference between the two signals is huge. The signal after the norm filter is much better then the signal prior to the norm filter. When looking at the F1 and accuracy scores of the two signals, the norm filter have only improved the signal with approximately 2%.

This raises the question: *is the F1 score a good indicator, and how high does the score need to be before it is application ready?* The F1 score seems to be generality adopted due to its simplicity and and acceptance in other areas of binary classification [27]. The F1 score required for a acceptable application can be dependent on many aspects. During the experiments, applications with a F1 score lower than 0.8 was generality influence much by interference from other devices. Where F1 scores higher than 0.8 could be filtered more easily, and gave more intuitive correct results.

The conclusion being that a functioning NILM solution requires commitment. A successful system will require the metering company to develop, or at least install, a meter capable of high speed data collection or processing. Furthermore must the residents have a interest in helping improving the accuracy, in order to get the best results. This could be done by registering appliances on a smart phone or tablet as suggested by Weiss [25].

7.4 NILM Privacy Concerns

As shown in Chapter 6 it is possible to gesture about the TV habits of a household. Potentially can the usage pattern for any appliance be derived using NILM techniques. Researchers have shown that even the current movie shown on a TV can be deduced by looking at the power profile [?]. With NILM techniques it is more or less possible to track the behaviour of a residential household [?].

This kind of possible surveillance raises some privacy concerns. Some simple techniques can be used to mask the appliances from a NILM applications, to improve privacy. It have been shown that down-sampling the consumption signal makes it harder for the NILM algorithms. Quantifying the signal by rounding it to a multiple of a pre-defined quantization factor also helps mask signals. Averaging the consumption signal over a time period is also shown to have a masking effect [?]. Common for all these techniques is that they must be actively build in the meter and transfer mechanisms to mask the appliance data. This will only be the case if the meter company have the same privacy concerns as the users.

An other approach is to use non-intrusive load leveling (NILL) to mask the appliance signals. NILL is a technique where a capacitive device like a battery is used to store energy. This energy is used to camouflage power signatures from the household [?].

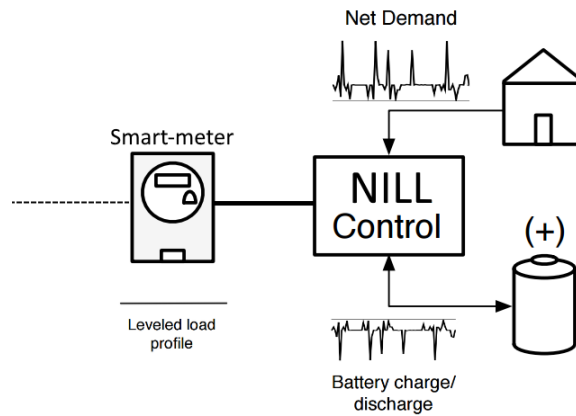


Figure 7.2: NILL illustration. Inspired by [?]

This principle is illustrated in figure 7.2. A NILL solution can be installed after the smart-meter and can provide privacy, without the consent of the electricity utility providers.

At the current state of NILM is the technology not accurate enough to make reliable detailed surveillance of the household residents. But still advanced enough to raise some privacy concerns. There are many benefits with advancing the possibilities of NILM. Many services can be developed that can support and help the resident. Some of the possibilities are in saving systems and home automation. This systems could greatly benefit the resident of the household. Electronic companies could also benefit from getting accurate user statistics for there devices. But advancing the technology in this direction would also serve to expose the privacy of the home. This raises the quotations about who are allowed accesses such sensitive information, and for what purposes must it be used.

Conclusion 8

Bibliography

- [1] A. Zoha, A. Gluhak, M. A. Imran, and S. Rajasegarar, “Non-intrusive load monitoring approaches for disaggregated energy sensing: A survey,” *Sensors*, vol. 12, no. 12, pp. 16838–16866, 2012. Appliance Load Monitoring (ALM) is essential for energy management solutions, allowing them to obtain appliance-specific energy consumption statistics that... Appliance Load Monitoring (ALM) is essential for energy management solutions, allowing them to obtain appliance-specific energy consumption statistics that can...
- [2] Z. Ghahramani and M. I. Jordan, “Factorial hidden markov models,” *Machine Learning*, vol. 29, no. 2, pp. 245–273, 1997. Hidden Markov models (HMMs) have proven to be one of the most widely used tools for learning probabilistic models of time series data. In an HMM, information... Hidden Markov models (HMMs) have proven to be one of the most widely used tools for learning probabilistic models of time series data. In an HMM, information...
- [3] O. Parson, S. Ghosh, M. Weal, and A. Rogers, “Non-intrusive load monitoring using prior models of general appliance types,” in *Proceedings of the Twenty-Sixth Conference on Artificial Intelligence (AAAI-12)*, pp. 356–362, July 2012.
- [4] C. Beckel, W. Kleiminger, R. Cicchetti, T. Staake, and S. Santini, “The eco data set and the performance of non-intrusive load monitoring algorithms,” in *Proceedings of the 1st ACM International Conference on Embedded Systems for Energy-Efficient Buildings (BuildSys 2014)*. Memphis, TN, USA, pp. 80–89, ACM, nov 2014.
- [5] D. Bonino, F. Corno, and L. D. Russis, “Home energy consumption feedback: A user survey,” *Energy and Buildings*, vol. 47, pp. 383–393, 2012.
- [6] N. Regnauld, “Generalisation and data quality,” 2015-08-01T00:00:00Z. Type: article; CC BY.
- [7] I. 8402:1994, “Quality management and quality assurance – vocabulary,” 1994-03-24.
- [8] S. Kelling, D. Fink, F. A. L. Sorte, A. Johnston, N. E. Bruns, and W. M. Hochachka, “Taking a ‘big data’ approach to data quality in a citizen science project,” 2015-10-27; 2015-11. Type: Text; © The Author(s) 2015; Open AccessThis article is distributed under the terms of the Creative Commons Attribution 4.0 International License (<http://creativecommons.org/licenses/by/4.0/>), which permits unrestricted use, distribution, and reproduction in any medium, provided you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons license, and indicate if changes were made.
- [9] I. 19157:2013, “Geographic information – data quality,” 2013-12-15.

- [10] G. Pastorello, D. Agarwal, T. Samak, C. Poindexter, B. Faybishenko, D. Gunter, R. Hollowgrass, D. Papale, C. Trotta, A. Ribeca, and E. Canfora, "Observational data patterns for time series data quality assessment," in *e-Science (e-Science), 2014 IEEE 10th International Conference on*, vol. 1, pp. 271–278, 2014. ID: 1.
- [11] M. Meijer, L. A. E. Vullings, J. D. Bulens, F. I. Rip, M. Boss, G. Hazeu, and M. Storm, "Spatial data quality and a workflow tool," 2015-08-01T00:00:00Z. Type: article; CC BY.
- [12] A. Simader, B. Kluger, N. Neumann, C. Bueschl, M. Lemmens, G. Lirk, R. Krska, and R. Schuhmacher, "Qcscreen: a software tool for data quality control in lc-hrms based metabolomics," 2015-10-24. Type: Software; Copyright 2015 Simader et al.
- [13] N. Batra, J. Kelly, O. Parson, H. Dutta, W. Knottenbelt, A. Rogers, A. Singh, and M. Srivastava, "Nilmk: An open source toolkit for non-intrusive load monitoring," 2014-04-15. Type: text.
- [14] D. G. Manolakis and V. K. Ingle, *Applied digital signal processing : theory and practice*. New York: Cambridge University Press, 2011.
- [15] P. J. S. G. Ferreira, "Interpolation and the discrete papoulis-gerchberg algorithm," *IEEE Transactions on Signal Processing*, vol. 42, no. 10, pp. 2596–2606, 1994. Analyze the performance of an iterative algorithm, similar to the discrete Papoulis-Gerchberg algorithm, and which can be used to recover missing samples in... An analysis of the performance of an iterative algorithm that can be used to recover missing samples in finite-length records of band-limited data.No...
- [16] M. Marques, A. Neves, J. Marques, and J. Sanches, "The papoulis-gerchberg algorithm with unknown signal bandwidth," vol. 4141, pp. 436–445, 2006.
- [17] D. J. Thomson, L. J. Lanzerotti, and C. G. Macleannan, "Interplanetary magnetic field: Statistical properties and discrete modes," *Journal of Geophysical Research*, vol. 106, no. A8, pp. 15941–15962, 2001.
- [18] D. Kondrashov and M. Ghil, "Spatio-temporal filling of missing points in geophysical data sets," *Nonlinear Processes in Geophysics*, vol. 13, no. 2, pp. 151–159, 2006. The majority of data sets in the geosciences are obtained from observations and measurements of natural systems, rather than in the laboratory. These data... The majority of data sets in the geosciences are obtained from observations and measurements of natural systems, rather than in the laboratory. These data sets...
- [19] A. Moghtaderi, P. Borgnat, and P. Flandrin, "Gap-filling by the empirical mode decomposition," 2012. We propose a novel gap-filling technique, based on the empirical mode decomposition (EMD). The idea is that a signal with missing data can be decomposed into a...
- [20] D. Srinivasan, W. S. Ng, and A. C. Liew, "Neural-network-based signature recognition for harmonic source identification," *IEEE Transactions on Power Delivery*, vol. 21, no. 1, pp. 398–405, 2006. This paper proposes a neural-network (NN)-based approach to nonintrusive harmonic source identification. In this approach, NNs are trained to extract important... Several NN-based classification models including multilayer perceptron (MLP), radial basis function (RBF) network, and support vector machines (SVM) with...

- [21] S. N. Patel, T. Robertson, J. A. Kientz, M. S. Reynolds, and G. D. Abowd, “At the flick of a switch: Detecting and classifying unique electrical events on,” in *the Residential Power Line*,” *Proc. 9th Int’l Conf. Ubiquitous Computing (UbiComp 07)*, ACM, pp. 271–288, Press, 2007.
- [22] J. Kelly and W. Knottenbelt, “Neural nilm: Deep neural networks applied to energy disaggregation,” 2015-07-23; 2015-09-28. Type: text.
- [23] R. Bonfigli, S. Squartini, M. Fagiani, and F. Piazza, “Unsupervised algorithms for non-intrusive load monitoring: An up-to-date overview,” 2015. Research on Smart Grids has recently focused on the energy monitoring issue, with the objective to maximize the user consumption awareness in building contexts...
- [24] J. Z. Kolter and T. Jaakkola, “Approximate inference in additive factorial hmms with application to energy disaggregation,” in *Proceedings of the Fifteenth International Conference on Artificial Intelligence and Statistics (AISTATS-12)* (N. D. Lawrence and M. A. Girolami, eds.), 2012.
- [25] M. Weiss, A. Helfenstein, F. Mattern, and T. Staake, “Leveraging smart meter data to recognize home appliances,” in *Pervasive Computing and Communications (PerCom), 2012 IEEE International Conference on*, pp. 190–197, 2012. ID: 1.
- [26] W. Kleiminger, C. Beckel, and S. Santini, “Household occupancy monitoring using electricity meters,” in *Proceedings of the 2015 ACM International Joint Conference on Pervasive and Ubiquitous Computing (UbiComp 2015)*, (Osaka, Japan), sep 2015.
- [27] S. Makonin and F. Popowich, “Nonintrusive load monitoring (nilm) performance evaluation,” *Energy Efficiency*, vol. 8, no. 4, pp. 809–814, 2014.
- [28] A. H. Kazmi, M. J. O’Grady, and G. M. O. Hare, *Towards Low-Cost Energy Monitoring*, pp. 2965–2970. Hershey, PA, USA: IGI Global, 2015. ID: 112719.